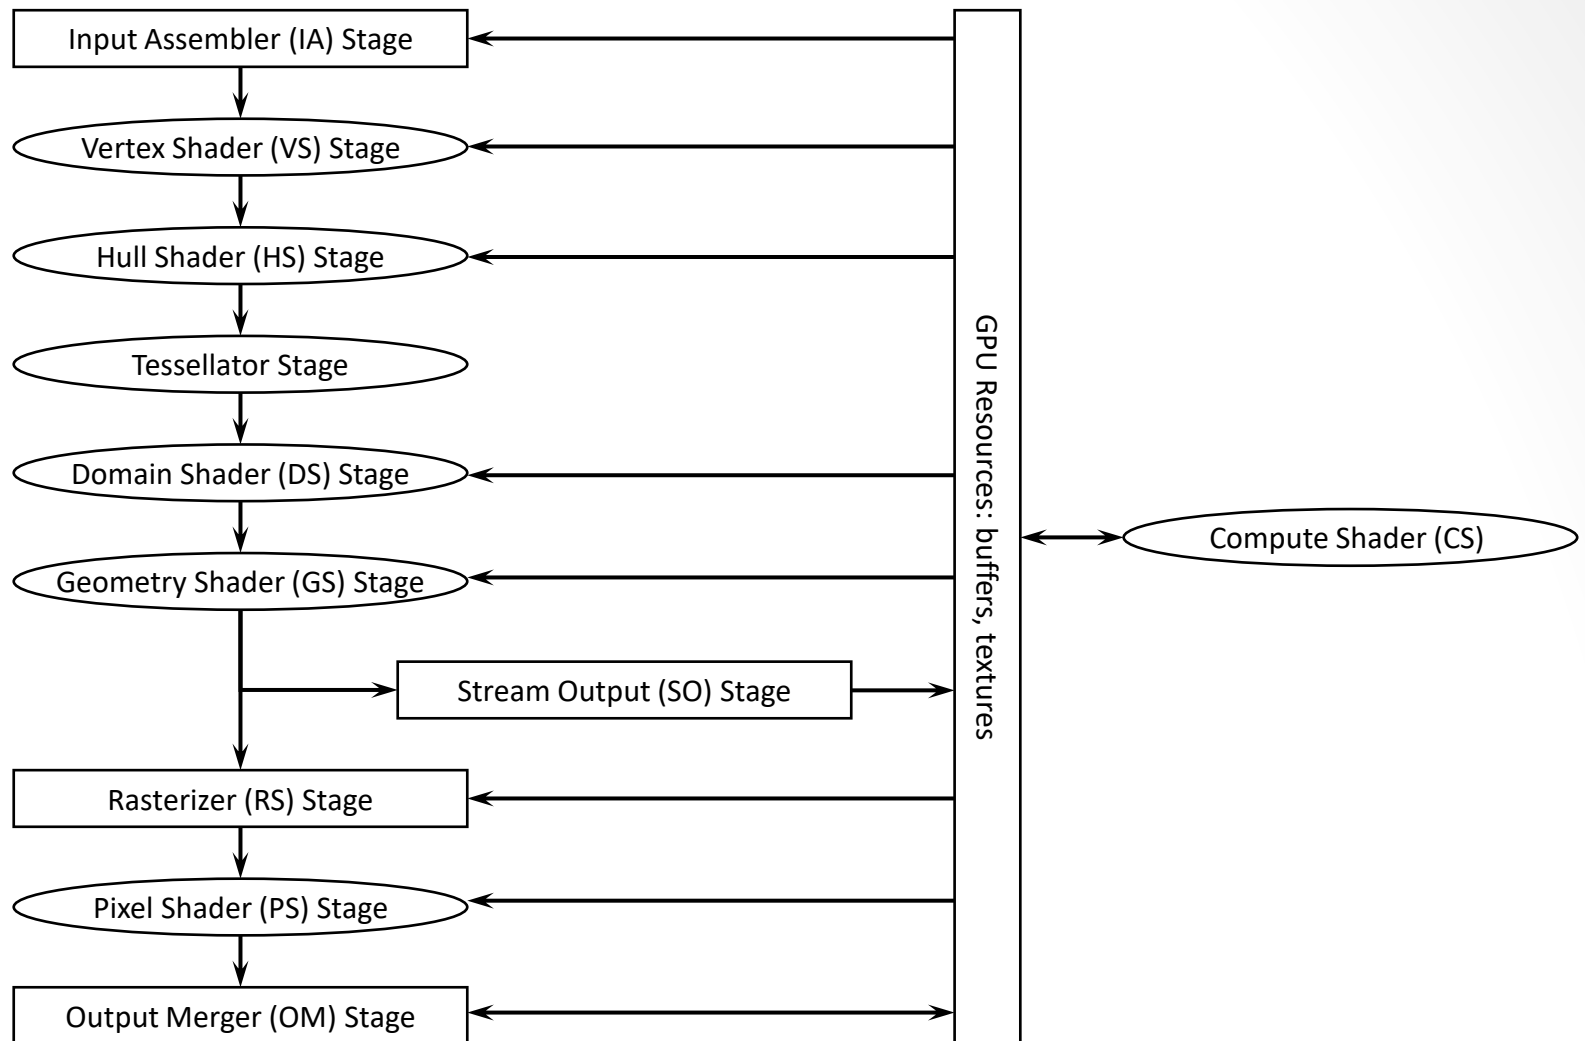


II. Direct3D Foundations

13. The Compute Shader

Game Graphic Programming
Kyung Hee University
Software Convergence
Prof. Daeho Lee

Compute Shader (1)



Compute Shader (2)

- The compute shader is not part of the rendering pipeline but sits off to the side.
- The compute shader can read and write to GPU resources.
- The compute shader can be mixed with graphics rendering, or used alone for GPGPU programming.

Threads & Thread Groups

- A thread group is executed on a single streaming multiprocessor.
 - GTX 1080: $20 \text{ SMs} \times \text{two blocks of } 64 \text{ cores} = 20 \times 2 \times 64 = 2,560 \text{ cores}$
 - RTX 4090: $128 \text{ SMs} = 128 \times 128 = 16,384 \text{ cores}$ (SM can contain up to 8 blocks)
 - A thread group consists of n threads.
- In Direct3D, thread groups are launched via the following method call:
 - `void ID3D12GraphicsCommandList::Dispatch (`
 `UINT ThreadGroupCountX,`
 `UINT ThreadGroupCountY,`
 `UINT ThreadGroupCountZ) ;`
 - This enables you to launch a 3D grid of thread groups.

The Components of CS (1)

- A compute shader consists of the following components:
 - 1. Global variable access via constant buffers.
 - 2. Input and output resources, which are discussed in the next section.
 - 3. The **[numthreads (X, Y, Z)]** attribute, which specifies the number of threads in the thread group as a 3D grid of threads.
 - 4. The shader body that has the instructions to execute for each thread.
 - 5. Thread identification system value parameters.
 - Each thread group is assigned an ID by the system: **SV_GroupID**.
 - Inside a thread group, each thread is given a unique ID relative to its group: **SV_GroupThreadID**.
 - A **Dispatch** call dispatches a grid of thread groups. The dispatch thread ID uniquely identifies a thread relative to all the threads generated by a **Dispatch** call: **SV_DispatchThreadID**.

The Components of CS (2)

```
// VecAdd.hlsl
struct Data
{
    float3 v1;
    float2 v2;
};

StructuredBuffer<Data> gInputA : register(t0);
StructuredBuffer<Data> gInputB : register(t1);
RWStructuredBuffer<Data> gOutput : register(u0);
    // unordered access view

[numthreads(32, 1, 1)]
void CS(int3 dtid : SV_DispatchThreadID)
{
    gOutput[dtid.x].v1 = gInputA[dtid.x].v1 + gInputB[dtid.x].v1;
    gOutput[dtid.x].v2 = gInputA[dtid.x].v2 + gInputB[dtid.x].v2;
}
```

The Components of CS (2)

```
// For textures
```

```
Texture2D gInputA;
```

```
Texture2D gInputB;
```

```
RWTexture2D<float4> gOutput;
```

```
[numthreads(16, 16, 1)]
```

```
void CS(int3 dispatchThreadID : SV_DispatchThreadID)
```

```
{
```

```
    gOutput[dispatchThreadID.xy] =
```

```
        gInputA[dispatchThreadID.xy] + gInputB[dispatchThreadID.xy];
```

```
}
```

Compile Shader & Compute PSO (1)

D3D12_COMPUTE_PIPELINE_STATE_DESC

```
typedef struct D3D12_COMPUTE_PIPELINE_STATE_DESC {  
    ID3D12RootSignature      *pRootSignature;  
    D3D12_SHADER_BYTECODE    CS;  
    UINT                      NodeMask;  
    D3D12_CACHED_PIPELINE_STATE CachedPSO;  
    D3D12_PIPELINE_STATE_FLAGS Flags;  
} D3D12_COMPUTE_PIPELINE_STATE_DESC;
```

- This describes a compute pipeline state object.

Compile Shader & Compute PSO (1)

```
void VecAddCSApp::BuildShadersAndInputLayout() {
    mShaders["vecAddCS"]
        = d3dUtil::CompileShader(L"Shaders\\VecAdd.hlsl", nullptr,
            "CS", "cs_5_0");
}

void VecAddCSApp::BuildPSOs() {
    D3D12_COMPUTE_PIPELINE_STATE_DESC computePsoDesc = {};
    computePsoDesc.pRootSignature = mRootSignature.Get();
    computePsoDesc.CS = {
        reinterpret_cast<BYTE*>(mShaders["vecAddCS"]->
            GetBufferPointer()),
        mShaders["vecAddCS"]->GetBufferSize()
    };
    computePsoDesc.Flags = D3D12_PIPELINE_STATE_FLAG_NONE;
    ThrowIfFailed(md3dDevice->
        CreateComputePipelineState(&computePsoDesc,
            IID_PPV_ARGS(&mPSOs["vecAdd"])));
}
```

Data Input & Output Resources (1)

- Texture inputs
 - Inputs are bound to shader resource view (SRVs).
- Texture outputs and unordered access views (UAVs)
 - Outputs are treated special and have the special prefix to their type "**RW**," which stands for read-write.
 - Inputs are bound to unordered access views (UAVs).

Data Input & Output Resources (2)

```

• void ID3D12Device::CreateUnorderedAccessView(
    ID3D12Resource          *pResource,
    ID3D12Resource          *pCounterResource,
    const D3D12_UNORDERED_ACCESS_VIEW_DESC *pDesc,
    [in] D3D12_CPU_DESCRIPTOR_HANDLE DestDescriptor
);

• typedef struct D3D12_UNORDERED_ACCESS_VIEW_DESC {
    DXGI_FORMAT          Format;
    D3D12_UAV_DIMENSION ViewDimension;
    union {
        D3D12_BUFFER_UAV          Buffer;
        D3D12_TEX1D_UAV           Texture1D;
        D3D12_TEX1D_ARRAY_UAV     Texture1DArray;
        D3D12_TEX2D_UAV           Texture2D;
        D3D12_TEX2D_ARRAY_UAV     Texture2DArray;
        D3D12_TEX2DMS_UAV         Texture2DMS;
        D3D12_TEX2DMS_ARRAY_UAV   Texture2DMSArray;
        D3D12_TEX3D_UAV           Texture3D;
    };
} D3D12_UNORDERED_ACCESS_VIEW_DESC;

```

Data Input & Output Resources (3)

```

• HRESULT ID3D12Device::CreateCommittedResource (
    [in]      const D3D12_HEAP_PROPERTIES *pHeapProperties,
    [in]      D3D12_HEAP_FLAGS             HeapFlags,
    [in]      const D3D12_RESOURCE_DESC    *pDesc,
    [in]      D3D12_RESOURCE_STATES        InitialResourceState,
    // D3D12_RESOURCE_STATE_UNORDERED_ACCESS
    [in, optional] const D3D12_CLEAR_VALUE *pOptimizedClearValue,
    [in]          REFIID                    riidResource,
    [out, optional] void                    **ppvResource
);

```

Data Input & Output Resources (4)

```

ComPtr<ID3D12Resource> mInputBufferA = nullptr;
ComPtr<ID3D12Resource> mInputUploadBufferA = nullptr;
ComPtr<ID3D12Resource> mInputBufferB = nullptr;
ComPtr<ID3D12Resource> mInputUploadBufferB = nullptr;
ComPtr<ID3D12Resource> mOutputBuffer = nullptr;
ComPtr<ID3D12Resource> mReadBackBuffer = nullptr;
const int NumDataElements = 32;

void VecAddCSApp::BuildBuffers() {
    std::ofstream fout("inputs.txt");
    std::vector<Data> dataA(NumDataElements);
    std::vector<Data> dataB(NumDataElements);
    for(int i = 0; i < NumDataElements; ++i) {
        dataA[i].v1 = XMFLOAT3(i, i, i);          dataA[i].v2 = XMFLOAT2(i, 0);
        dataB[i].v1 = XMFLOAT3(-i, i, 0.0f);      dataB[i].v2 = XMFLOAT2(0, -i);

        fout << "A(" << dataA[i].v1.x << ", " << dataA[i].v1.y << ", "
            << dataA[i].v1.z << ", " << dataA[i].v2.x << ", " << dataA[i].v2.y
            << ")" << "\t";
        fout << "B(" << dataB[i].v1.x << ", " << dataB[i].v1.y << ", "
            << dataB[i].v1.z << ", " << dataB[i].v2.x << ", " << dataB[i].v2.y
            << ")" << std::endl;
    }
}

```

Data Input & Output Resources (5)

```
UINT64 byteSize = dataA.size()*sizeof(Data);  
  
// Create some buffers to be used as SRVs.  
mInputBufferA = d3dUtil::CreateDefaultBuffer(  
    md3dDevice.Get(),  
    mCommandList.Get(),  
    dataA.data(),  
    byteSize,  
    mInputUploadBufferA);  
  
mInputBufferB = d3dUtil::CreateDefaultBuffer(  
    md3dDevice.Get(),  
    mCommandList.Get(),  
    dataB.data(),  
    byteSize,  
    mInputUploadBufferB);
```

Data Input & Output Resources (6)

```
// Create the buffer that will be a UAV.
ThrowIfFailed(md3dDevice->CreateCommittedResource(
    &CD3DX12_HEAP_PROPERTIES(D3D12_HEAP_TYPE_DEFAULT),
    D3D12_HEAP_FLAG_NONE,
    &CD3DX12_RESOURCE_DESC::Buffer(byteSize,
        D3D12_RESOURCE_FLAG_ALLOW_UNORDERED_ACCESS),
    D3D12_RESOURCE_STATE_UNORDERED_ACCESS,
    nullptr,
    IID_PPV_ARGS(&mOutputBuffer)));

ThrowIfFailed(md3dDevice->CreateCommittedResource(
    &CD3DX12_HEAP_PROPERTIES(D3D12_HEAP_TYPE_READBACK),
    D3D12_HEAP_FLAG_NONE,
    &CD3DX12_RESOURCE_DESC::Buffer(byteSize),
    D3D12_RESOURCE_STATE_COPY_DEST,
    nullptr,
    IID_PPV_ARGS(&mReadBackBuffer)));
}
```

Root Signatures

```
void VecAddCSApp::BuildRootSignature() {
    CD3DX12_ROOT_PARAMETER slotRootParameter[3];

    slotRootParameter[0].InitAsShaderResourceView(0);
    slotRootParameter[1].InitAsShaderResourceView(1);
    slotRootParameter[2].InitAsUnorderedAccessView(0);

    CD3DX12_ROOT_SIGNATURE_DESC rootSigDesc(3, slotRootParameter,
        0, nullptr, D3D12_ROOT_SIGNATURE_FLAG_NONE);

    ComPtr<ID3DBlob> serializedRootSig = nullptr;
    ComPtr<ID3DBlob> errorBlob = nullptr;
    HRESULT hr = D3D12SerializeRootSignature(&rootSigDesc,
        D3D_ROOT_SIGNATURE_VERSION_1,
        serializedRootSig.GetAddressOf(), errorBlob.GetAddressOf());
    if(errorBlob != nullptr) { /* ... */      ThrowIfFailed(hr);

    ThrowIfFailed(md3dDevice->CreateRootSignature(
        0, serializedRootSig->GetBufferPointer(),
        serializedRootSig->GetBufferSize(),
        IID_PPV_ARGS(mRootSignature.GetAddressOf())));
}
```


Working CS (1)

```
void VecAddCSApp::DoComputeWork() {
    ThrowIfFailed(mDirectCmdListAlloc->Reset());

    ThrowIfFailed(mCommandList->Reset(mDirectCmdListAlloc.Get(),
        mPSOs["vecAdd"].Get()));

    mCommandList->SetComputeRootSignature(mRootSignature.Get());

    mCommandList->SetComputeRootShaderResourceView(0,
        mInputBufferA->GetGPUVirtualAddress());
    mCommandList->SetComputeRootShaderResourceView(1,
        mInputBufferB->GetGPUVirtualAddress());
    mCommandList->SetComputeRootUnorderedAccessView(2,
        mOutputBuffer->GetGPUVirtualAddress());

    mCommandList->Dispatch(1, 1, 1);    // [numthreads(32, 1, 1)]
    // mCommandList->Dispatch(2, 1, 1);    // [numthreads(16, 1, 1)]
}
```

Working CS (2)

```
mCommandList->ResourceBarrier(1,
    &CD3DX12_RESOURCE_BARRIER::Transition(mOutputBuffer.Get(),
        D3D12_RESOURCE_STATE_COMMON, D3D12_RESOURCE_STATE_COPY_SOURCE));

mCommandList->CopyResource(mReadBackBuffer.Get(),
    mOutputBuffer.Get());

mCommandList->ResourceBarrier(1,
    &CD3DX12_RESOURCE_BARRIER::Transition(mOutputBuffer.Get(),
        D3D12_RESOURCE_STATE_COPY_SOURCE, D3D12_RESOURCE_STATE_COMMON));

ThrowIfFailed(mCommandList->Close());

ID3D12CommandList* cmdsLists[] = { mCommandList.Get() };
mCommandQueue->ExecuteCommandLists(_countof(cmdsLists), cmdsLists);

FlushCommandQueue();
```

Working CS (3)

```
// Map the data so we can read it on CPU.
Data* mappedData = nullptr;
ThrowIfFailed(mReadBackBuffer->Map(0, nullptr,
    reinterpret_cast<void**>(&mappedData)));

std::ofstream fout("results.txt");
for(int i = 0; i < NumDataElements; ++i) {
    fout << "(" << mappedData[i].v1.x << ", "
        << mappedData[i].v1.y << ", " << mappedData[i].v1.z
        << ", " << mappedData[i].v2.x << ", " << mappedData[i].v2.y
        << ")" << std::endl;
}

mReadBackBuffer->Unmap(0, nullptr);
}
```

Working CS (4)

```
// inputs.txt
```

```
A(0, 0, 0, 0, 0)  B(0, 0, 0, 0, 0)
A(1, 1, 1, 1, 0)  B(-1, 1, 0, 0, -1)
A(2, 2, 2, 2, 0)  B(-2, 2, 0, 0, -2)
A(3, 3, 3, 3, 0)  B(-3, 3, 0, 0, -3)
A(4, 4, 4, 4, 0)  B(-4, 4, 0, 0, -4)
A(5, 5, 5, 5, 0)  B(-5, 5, 0, 0, -5)
A(6, 6, 6, 6, 0)  B(-6, 6, 0, 0, -6)
A(7, 7, 7, 7, 0)  B(-7, 7, 0, 0, -7)
A(8, 8, 8, 8, 0)  B(-8, 8, 0, 0, -8)
A(9, 9, 9, 9, 0)  B(-9, 9, 0, 0, -9)
```

```
// ...
```

```
A(26, 26, 26, 26, 0)  B(-26, 26, 0, 0, -26)
A(27, 27, 27, 27, 0)  B(-27, 27, 0, 0, -27)
A(28, 28, 28, 28, 0)  B(-28, 28, 0, 0, -28)
A(29, 29, 29, 29, 0)  B(-29, 29, 0, 0, -29)
A(30, 30, 30, 30, 0)  B(-30, 30, 0, 0, -30)
A(31, 31, 31, 31, 0)  B(-31, 31, 0, 0, -31)
```

Working CS (5)

```
// results.txt
```

```
(0, 0, 0, 0, 0)
```

```
(0, 2, 1, 1, -1)
```

```
(0, 4, 2, 2, -2)
```

```
(0, 6, 3, 3, -3)
```

```
(0, 8, 4, 4, -4)
```

```
(0, 10, 5, 5, -5)
```

```
(0, 12, 6, 6, -6)
```

```
(0, 14, 7, 7, -7)
```

```
(0, 16, 8, 8, -8)
```

```
(0, 18, 9, 9, -9)
```

```
// ...
```

```
(0, 52, 26, 26, -26)
```

```
(0, 54, 27, 27, -27)
```

```
(0, 56, 28, 28, -28)
```

```
(0, 58, 29, 29, -29)
```

```
(0, 60, 30, 30, -30)
```

```
(0, 62, 31, 31, -31)
```

Blur Example



Sobel Example

