

Operating System

Ch12: I/O systems

BeomSeok Kim

Department of Computer Engineering

KyungHee University

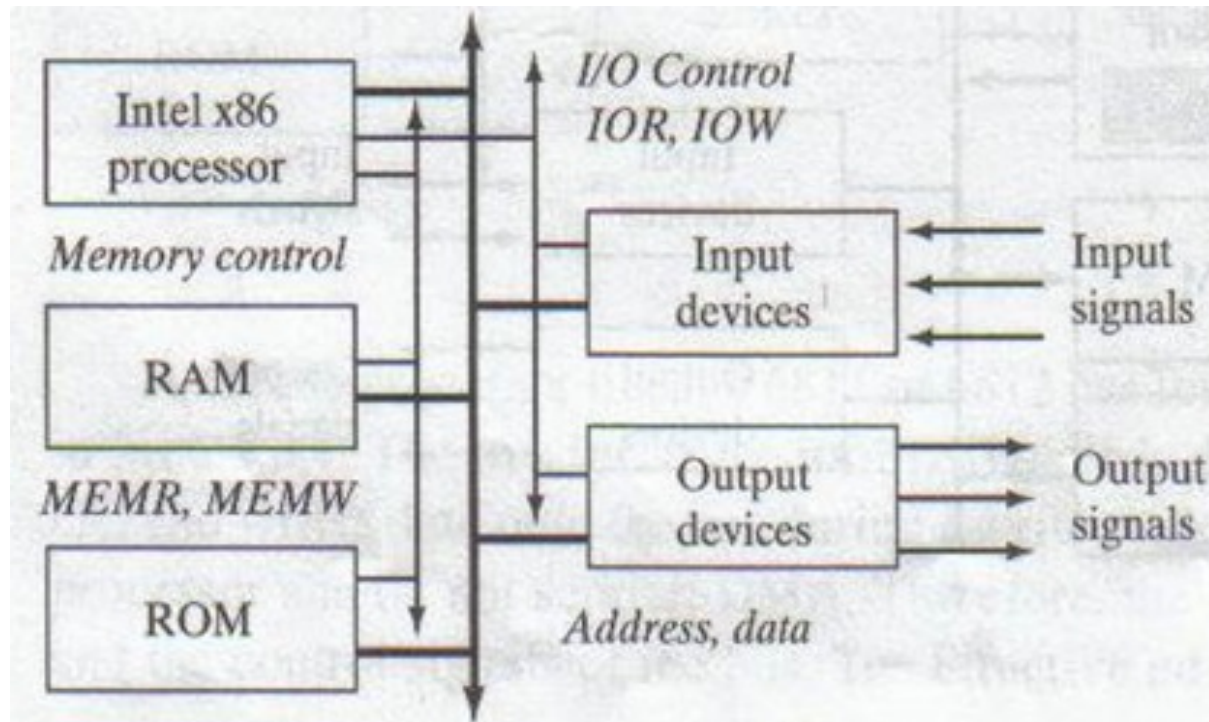
passion0822@khu.ac.kr

- I/O operation

- ✓ I/O request via I/O instruction
 - Direct I/O vs. Memory-mapped I/O
 - Communicates with registers in I/O controller
 - Typically, IR (Instruction Register) & DR (Data Register)
- ✓ I/O method
 - Programmed I/O
 - Interrupt
 - DMA (Direct Memory Access)

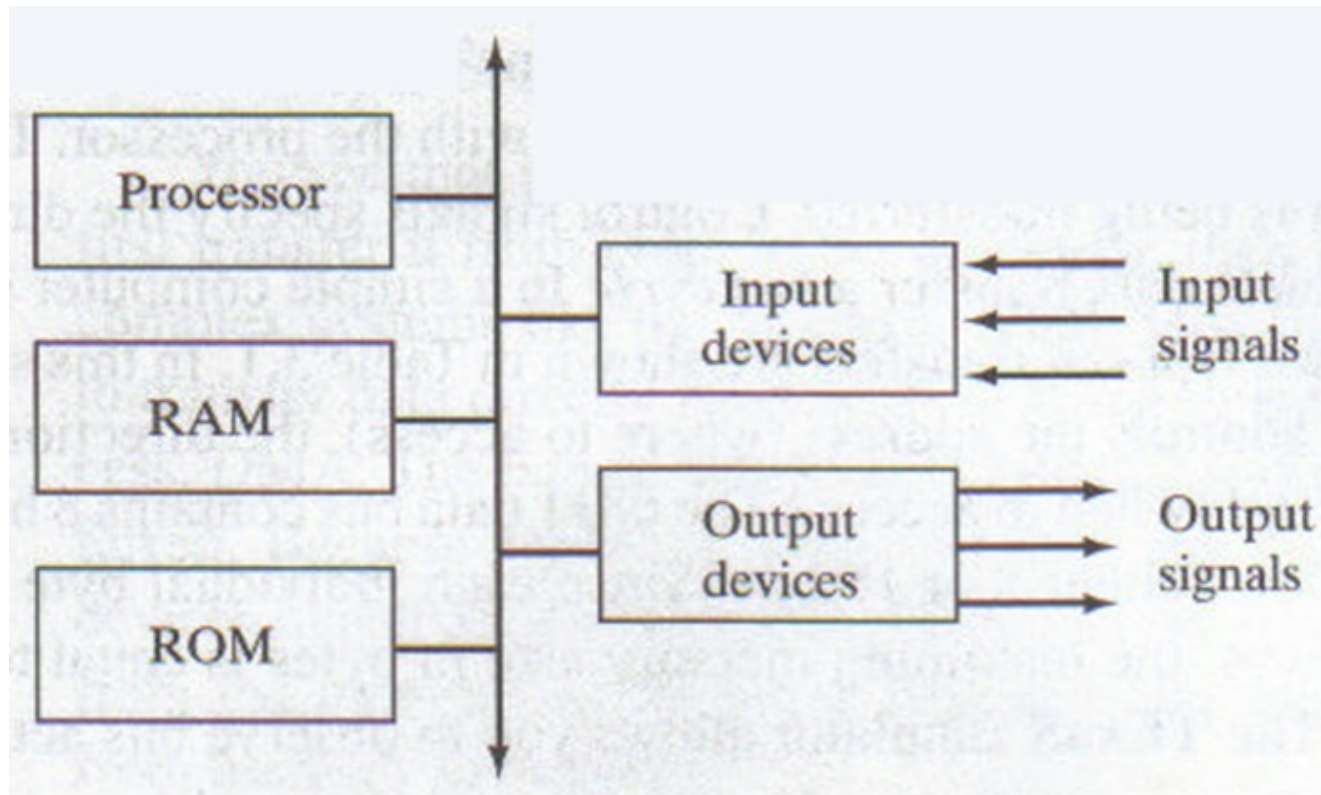
Direct I/O

- Separate addresses: memory vs. I/O
- Separate instructions
 - ✓ MEMR / MEMW (Load / Store)
 - ✓ IOR / IOW (IN / OUT)

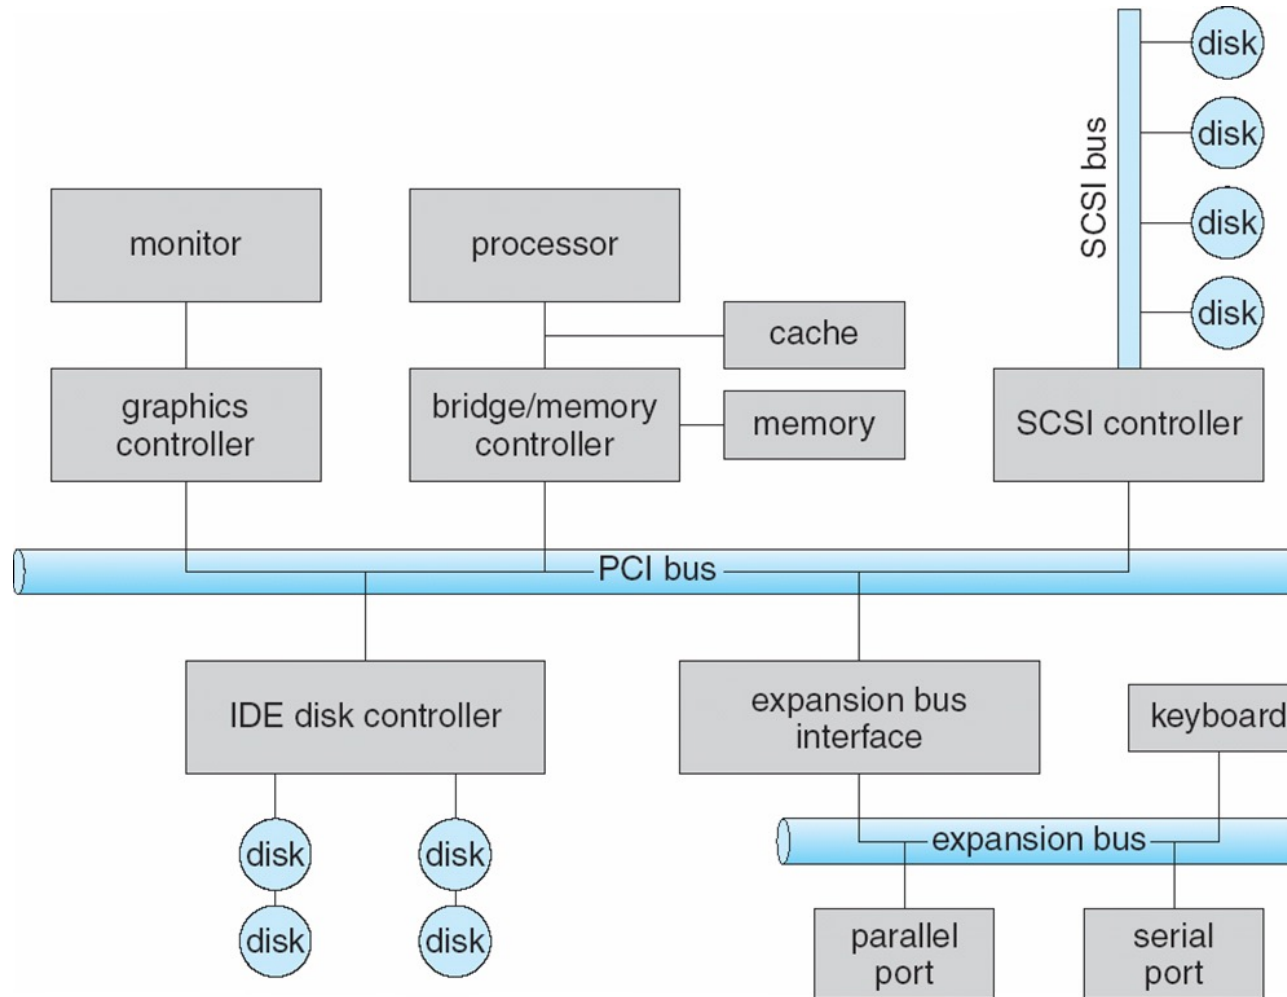


Memory-Mapped I/O

- Integrated addresses: memory & I/O
- Integrated instructions
 - ✓ MEMR / MEMW (Load / Store)



A Typical PC Bus Structure

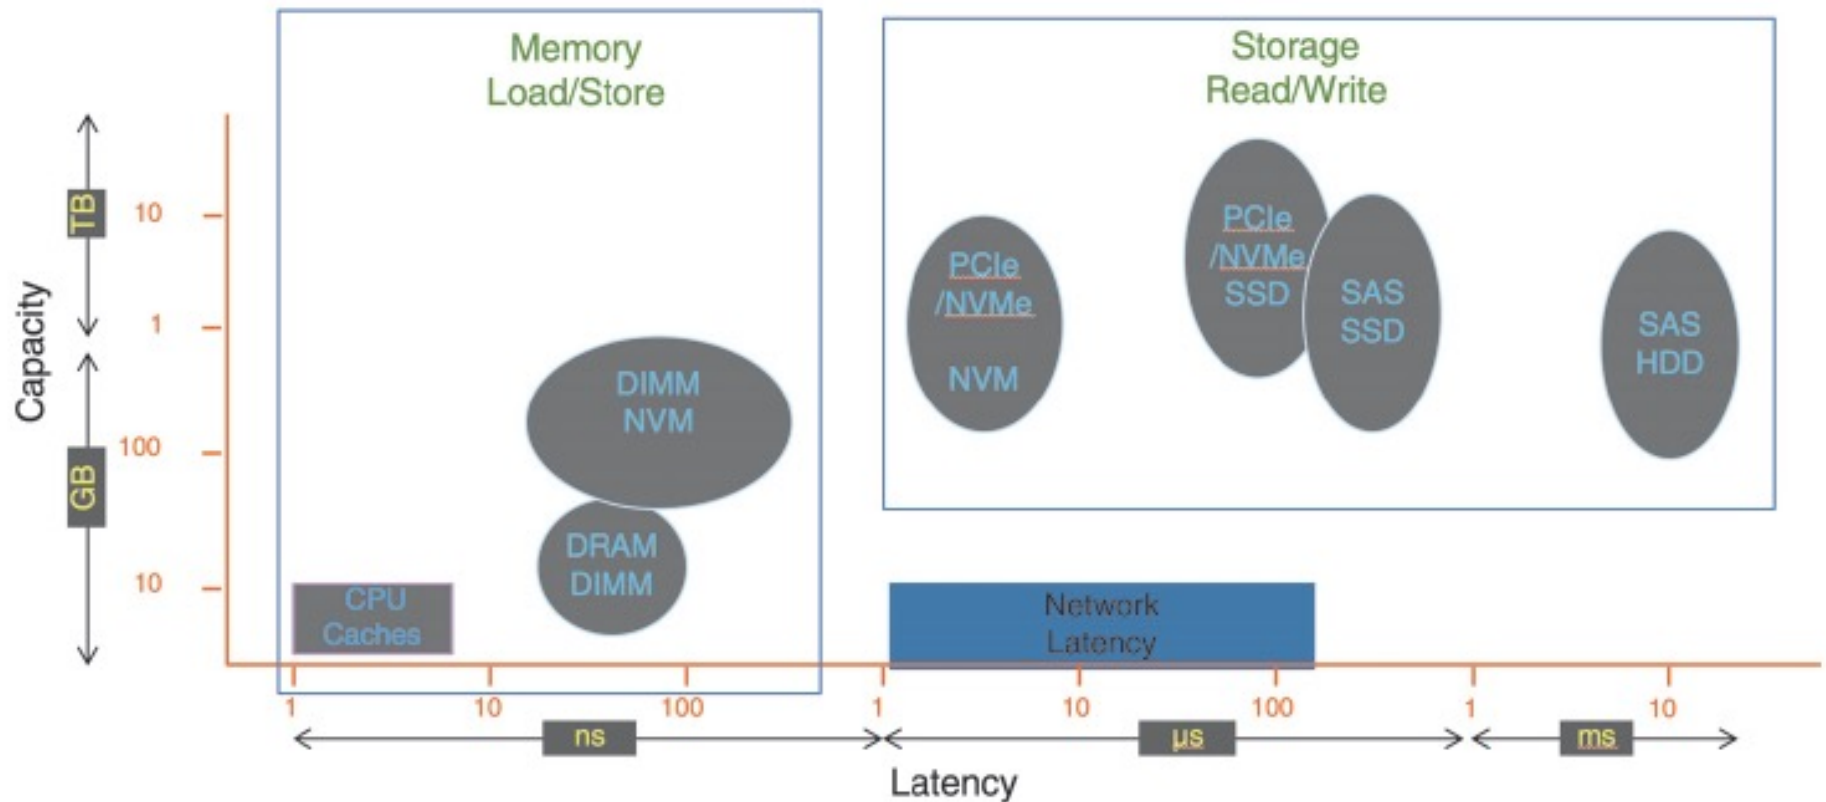


Device Controller

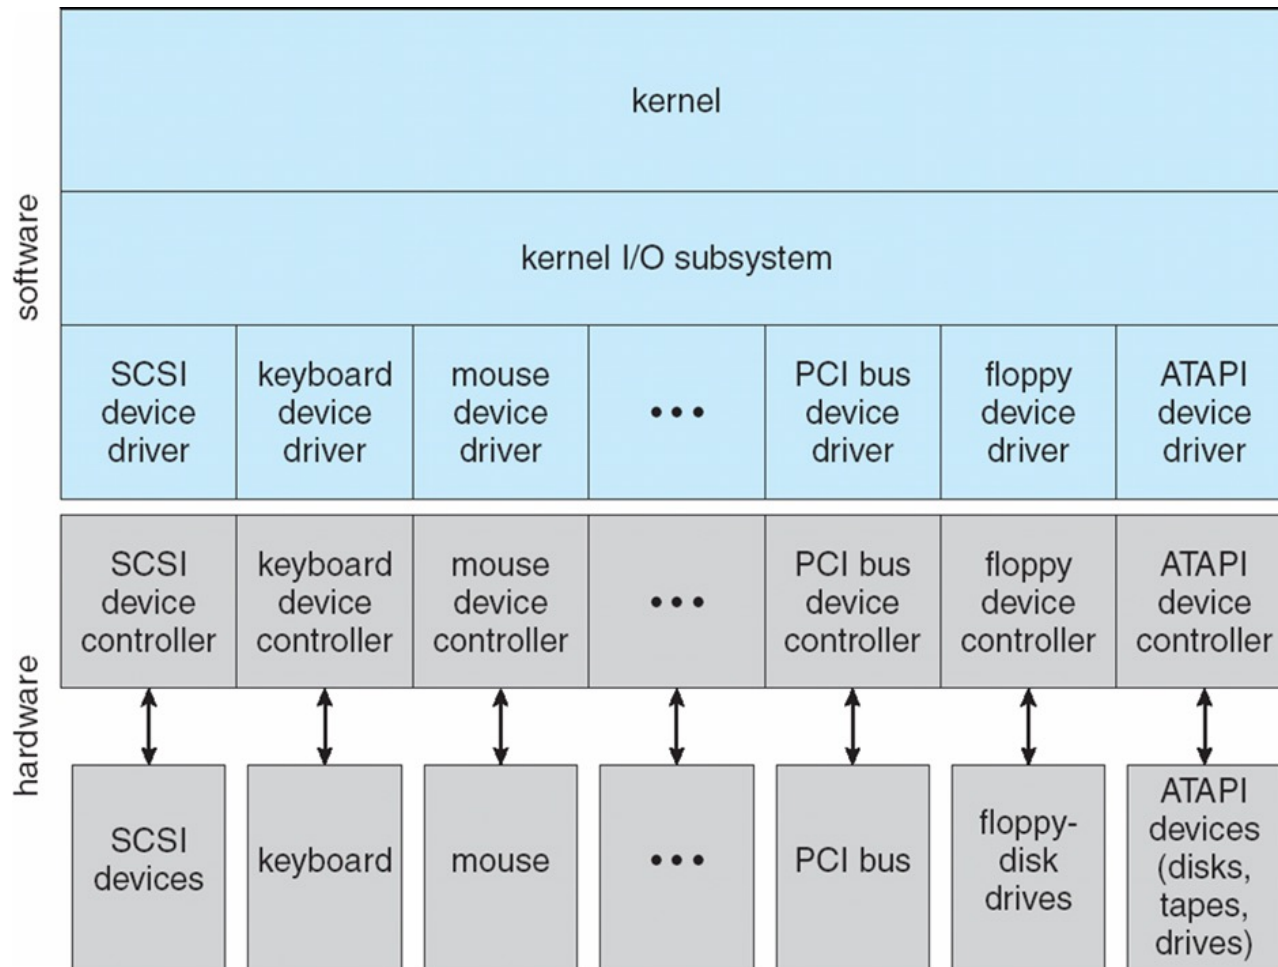


- Device controller (or host adapter)
 - ✓ I/O devices have components:
 - Mechanical component
 - Electronic component
 - ✓ The electronic component is the device controller
 - May be able to handle multiple devices
 - ✓ Controller's tasks
 - Convert serial bit stream to block of bytes
 - Perform error correction as necessary
 - Make available to main memory

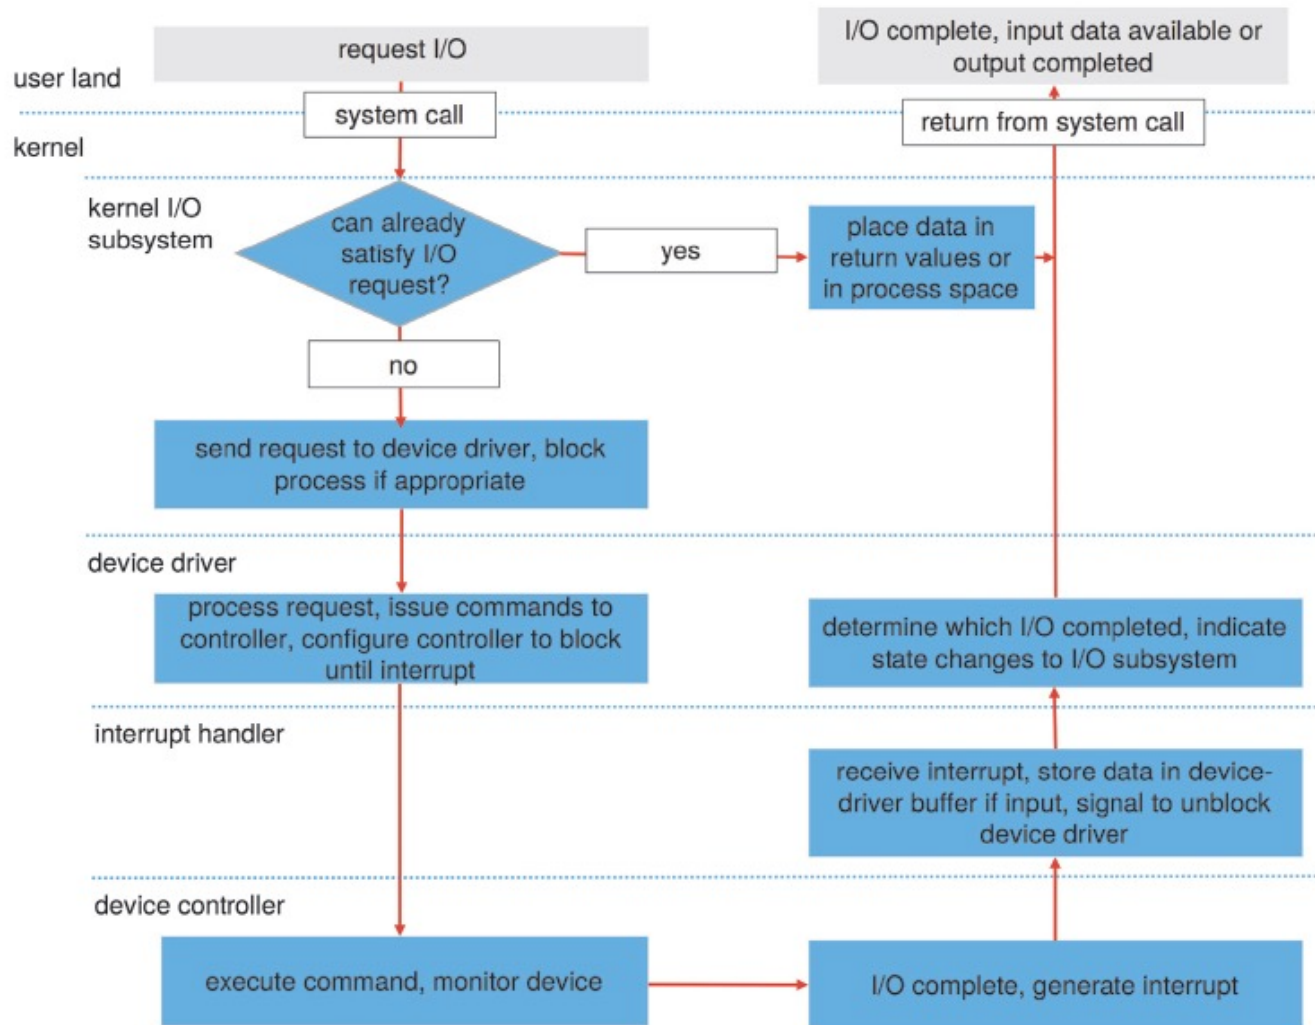
I/O Performance of Storage and Network Latency



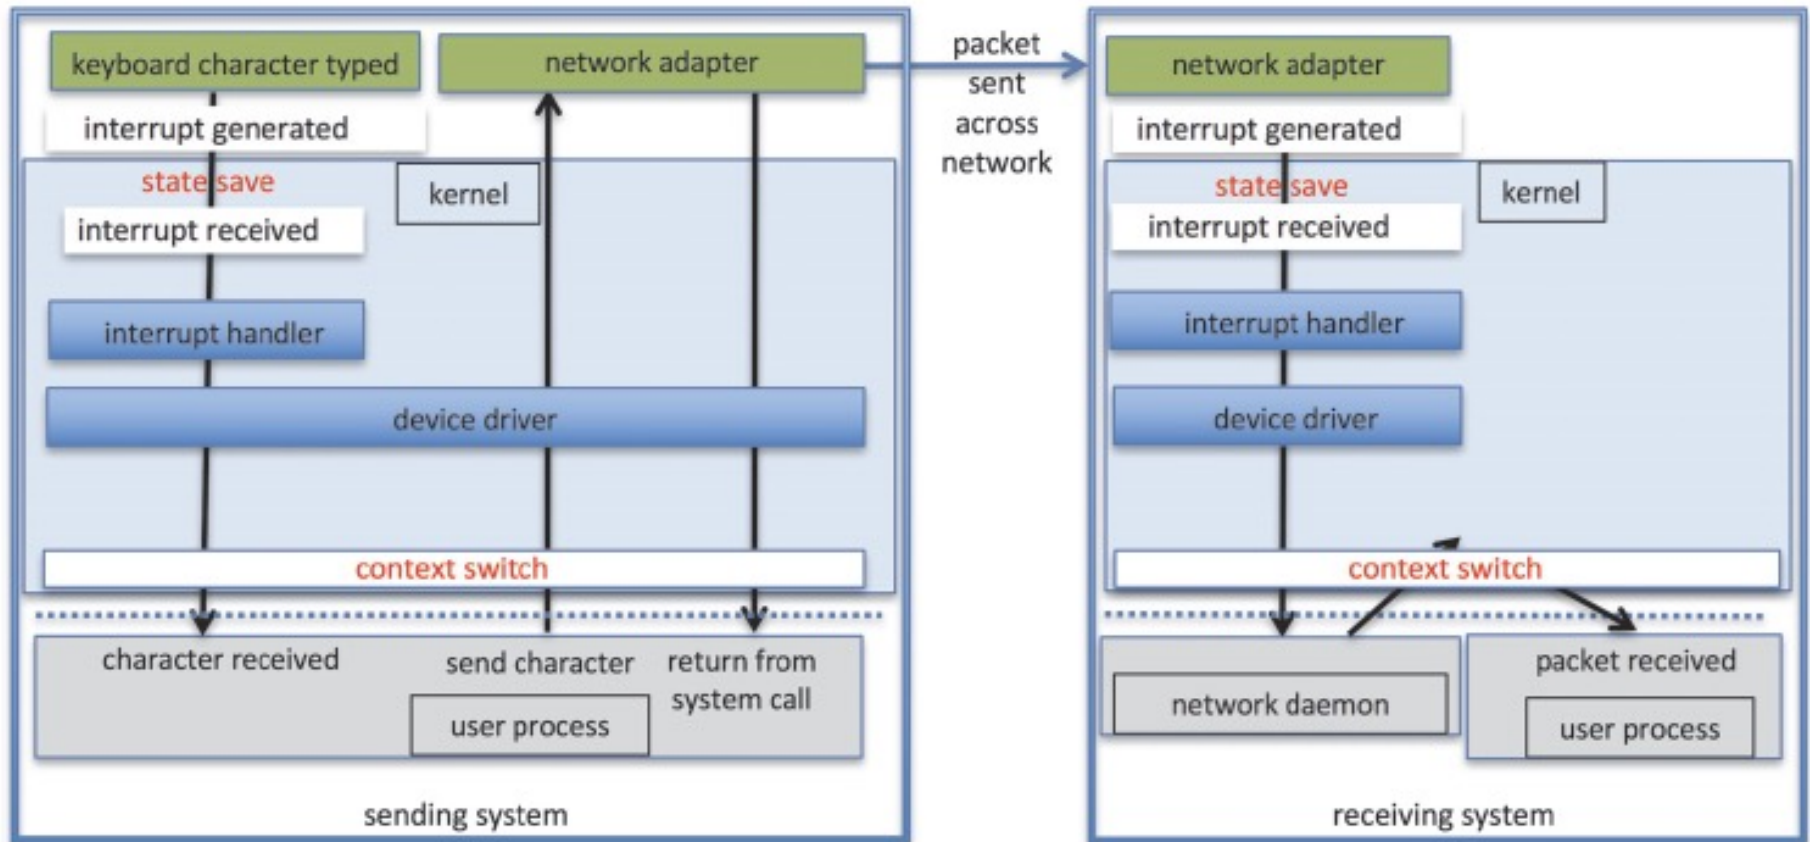
Kernel I/O Structure



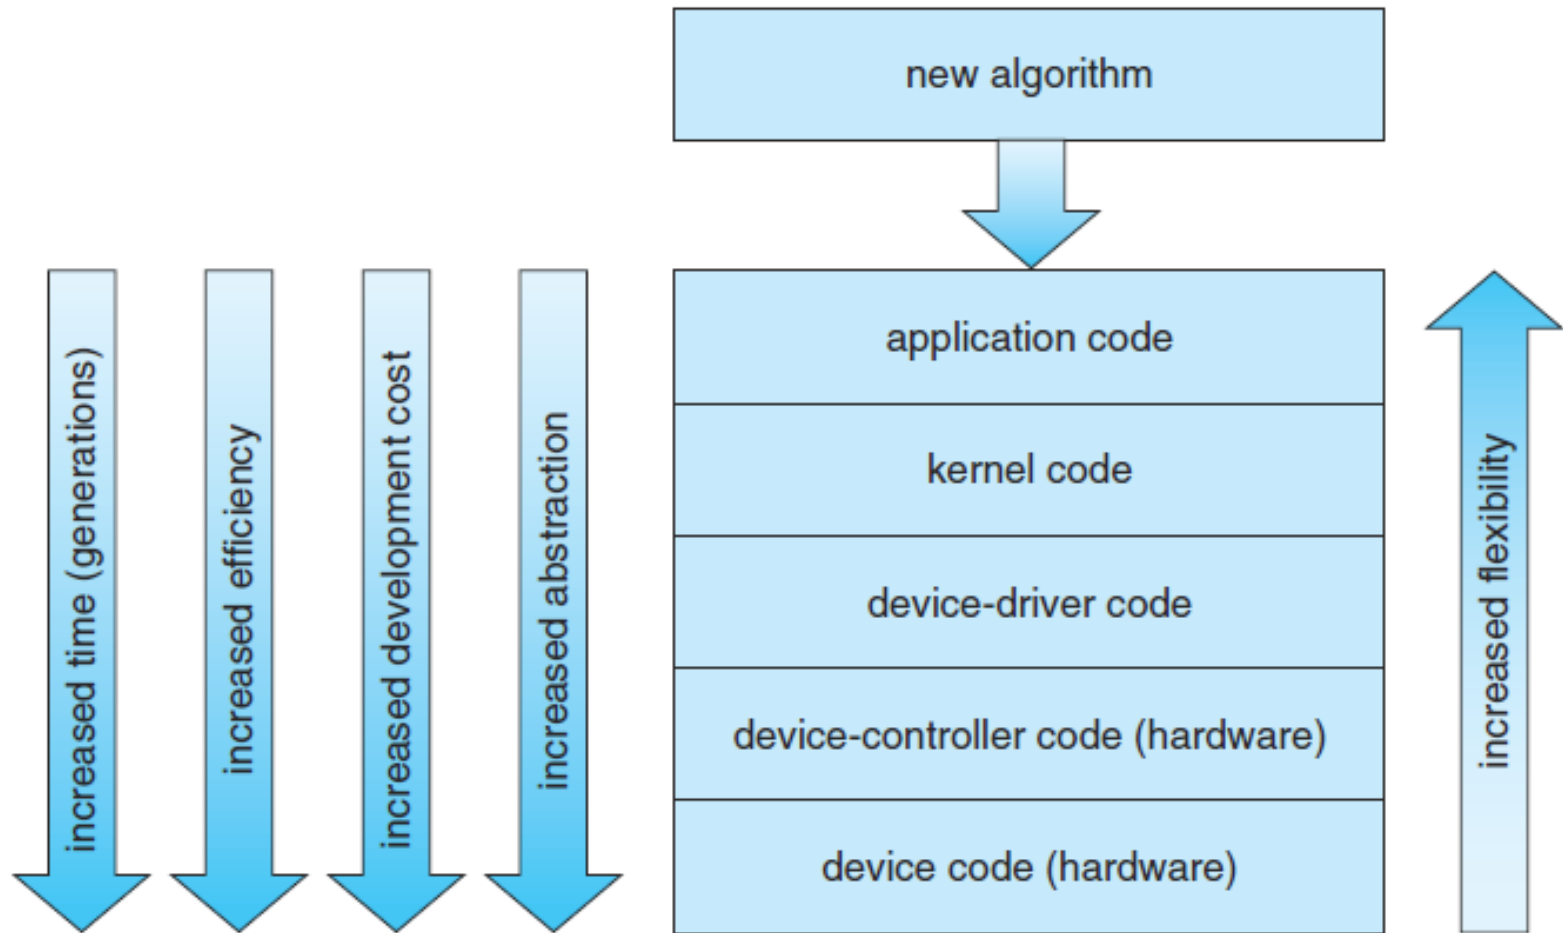
Life Cycle of An I/O Request



Intercomputer Communications



Device-Functionality Progression



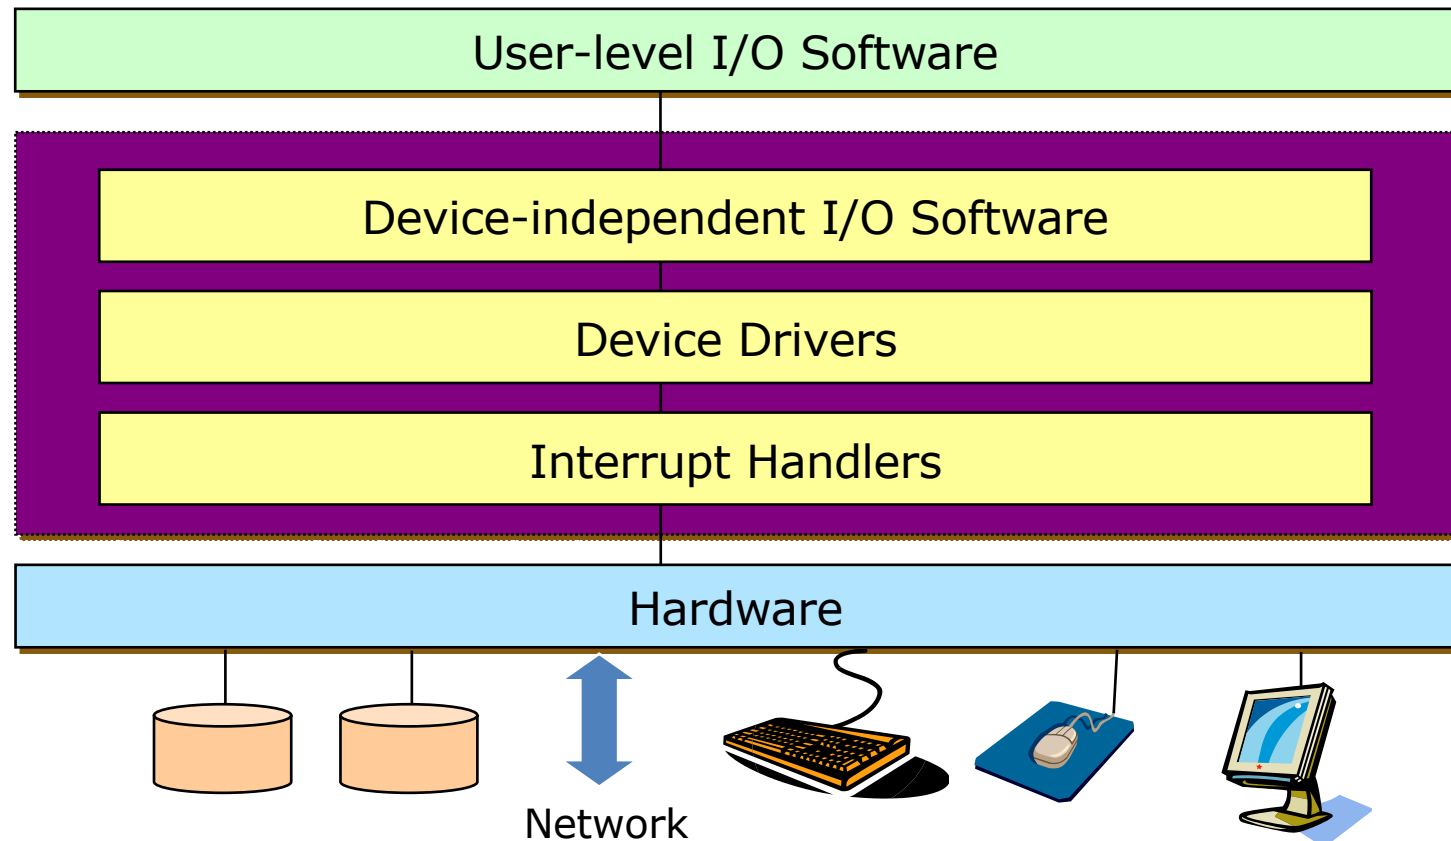
Goals of I/O Software



■ Goals

- ✓ Device independence
 - Programs can access any I/O device without specifying device in advance
- ✓ Uniform naming
 - Name of a file or device should simply be a string or an integer
- ✓ Error handling
 - Handle as close to the hardware as possible
- ✓ Synchronous vs. asynchronous
 - blocked transfers vs. interrupt-driven
- ✓ Buffering
 - Data coming off a device cannot be stored in final destination
- ✓ Sharable vs. dedicated devices
 - Disks vs. tape drives
 - Unsharable devices introduce problems such as deadlocks


I/O Software Layers



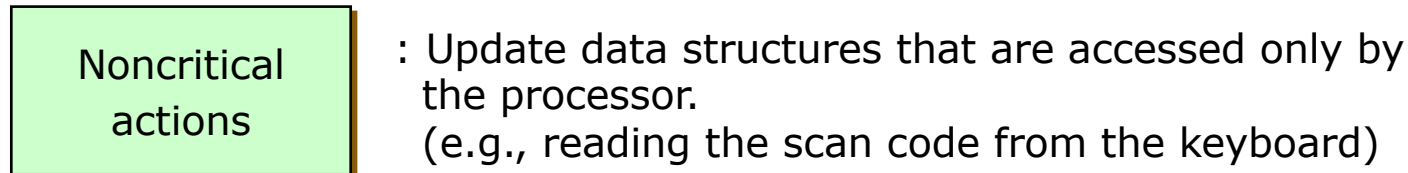
Interrupt Handlers



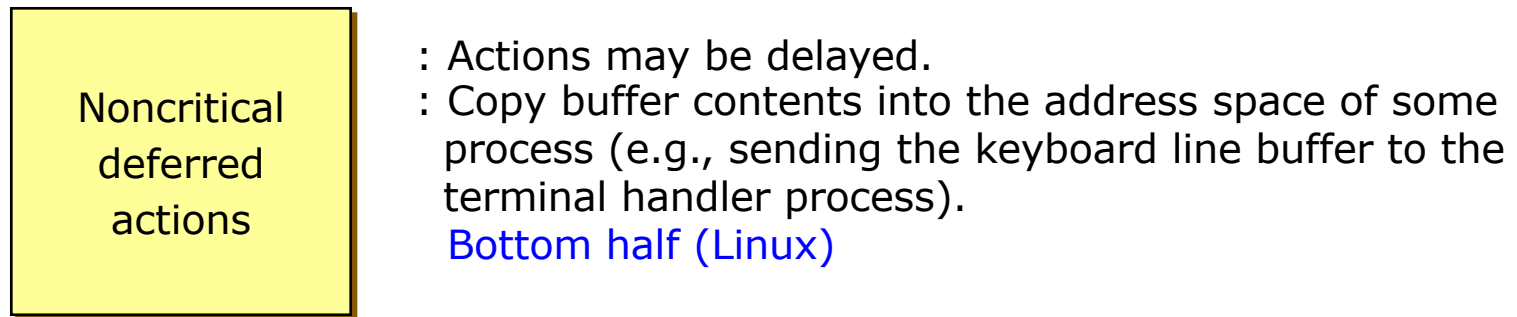
■ Handling interrupts



Reenable interrupts



Return from interrupts



Device Drivers



■ Device drivers

- ✓ Device-specific code to control each I/O device interacting with device-independent I/O software and interrupt handlers
- ✓ Requires to define a well-defined model and a standard interface of how they interact with the rest of the OS
- ✓ Implementing device drivers:
 - Statically linked with the kernel
 - Selectively loaded into the system during boot time
 - Dynamically loaded into the system during execution (especially for hot pluggable devices)

Device-Independent I/O Software

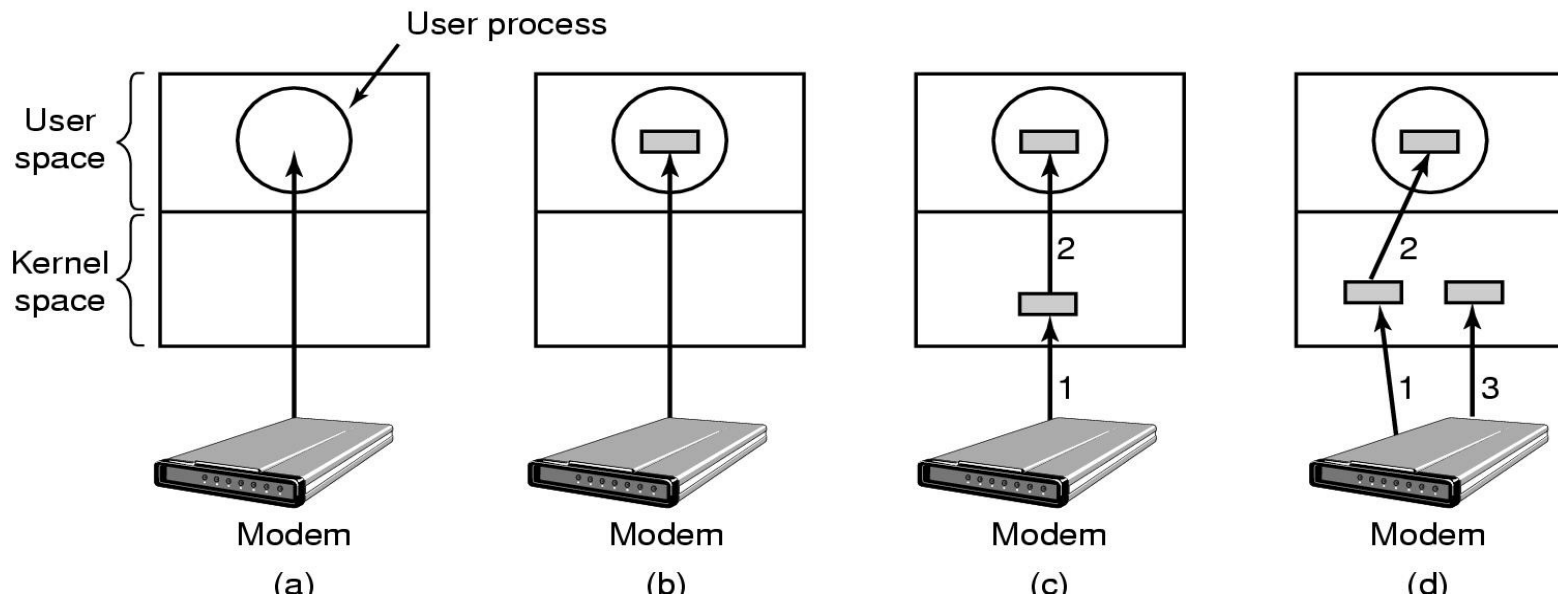


- Uniform interfacing for device drivers
 - ✓ In Unix, devices are modeled as special files
 - They are accessed through the use of system calls such as `open()`, `read()`, `write()`, `close()`, `ioctl()`, etc.
 - A file name is associated with each device
 - ✓ Major device number locates the appropriate driver
 - Minor device number (stored in i-node) is passed as a parameter to the driver in order to specify the unit to be read or written
 - ✓ The usual protection rules for files also apply to I/O devices

Device-Independent I/O Software

■ Buffering

- ✓ (a) Unbuffered
- ✓ (b) Buffered in user space
- ✓ (c) Buffered in the kernel space
- ✓ (d) Double buffering in the kernel



Device-Independent I/O Software



■ Error reporting

- ✓ Many errors are device-specific and must be handled by the appropriate driver, but the framework for error handling is device independent
- ✓ Programming errors vs. actual I/O errors
- ✓ Handling errors
 - Returning the system call with an error code
 - Retrying a certain number of times
 - Ignoring the error
 - Killing the calling process
 - Terminating the system

Device-Independent I/O Software



- Allocating and releasing dedicated devices
 - ✓ Some devices cannot be shared
 - (1) Require processes to perform `open()`'s on the special files for devices directly
 - The process retries if `open()` fails
 - (2) Have special mechanisms for requesting and releasing dedicated devices
 - An attempt to acquire a device that is not available blocks the caller

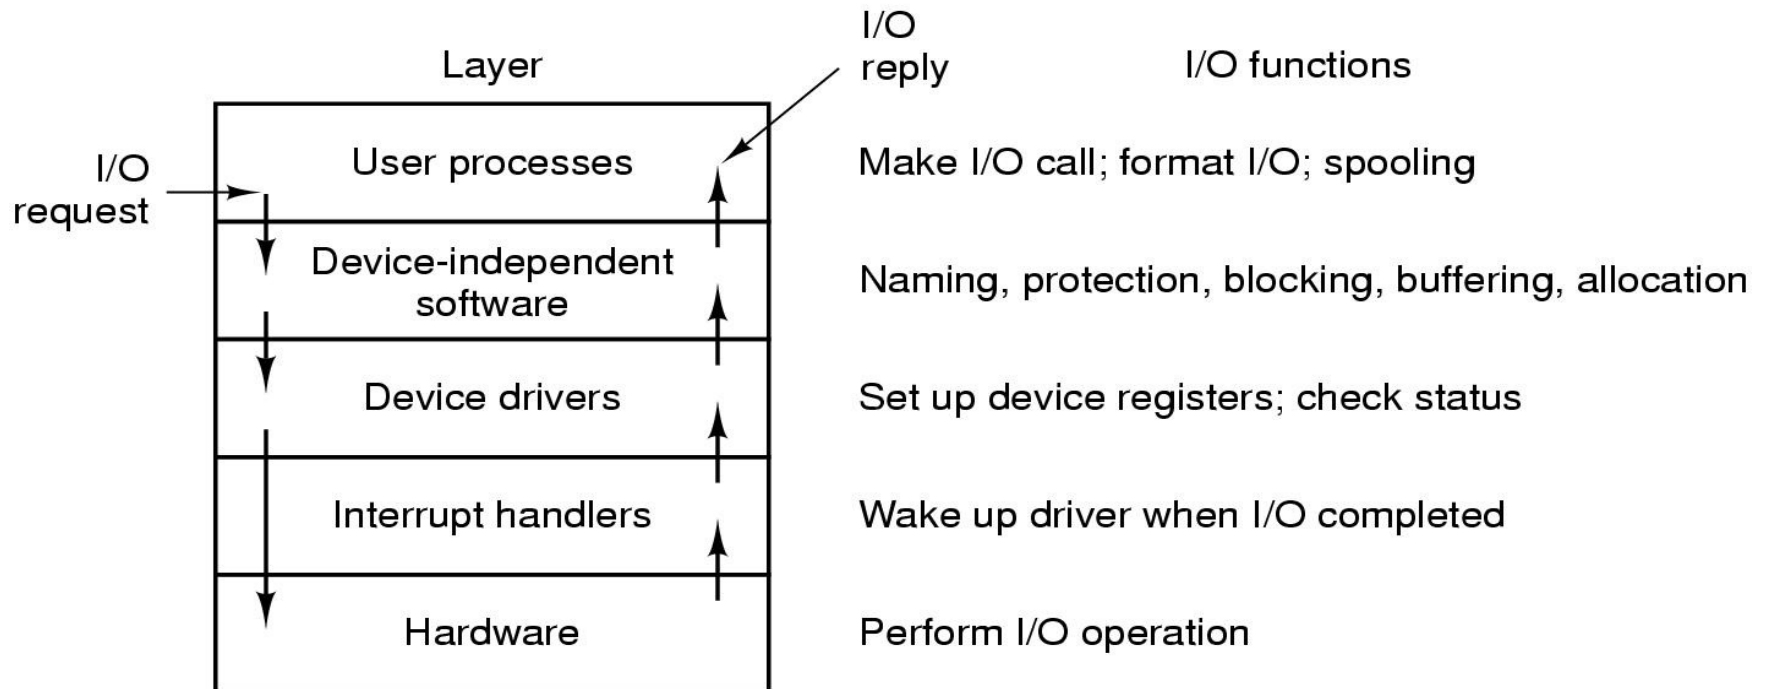
- Device-independent block size
 - ✓ Treat several sectors as a single logical block
 - ✓ The higher layers only deal with abstract devices that all use the same block size

User-Space I/O Software



- Provided as a library
 - ✓ Standard I/O library in C
 - `fopen()` vs. `open()`
- Spooling
 - ✓ A way of dealing with dedicated I/O devices in a multiprogramming system
 - ✓ Implemented by a daemon and a spooling directory
 - ✓ Printers, network file transfers, USENET news, mails, etc.

I/O Systems Layers



Thank You!
Q&A