

Operating System

Ch04: Thread & Concurrency

BeomSeok Kim

Department of Computer Engineering

KyungHee University

passion0822@khu.ac.kr

Process



■ Heavy-weight



- ✓ A process includes many things:
 - An address space (all the code and data pages)
 - OS resources (e.g., open files) and accounting info.
 - Hardware execution state (PC, SP, registers, etc.)
- ✓ Creating a new process is costly because all of the data structures must be allocated and initialized
 - Linux: over 100 fields in `task_struct`
(excluding page tables, etc.)
- ✓ Inter-process communication is costly, since it must usually go through the OS
 - Overhead of system calls and copying data

Thread Concept: Key Idea



- Separate the concept of a process from its execution state
 - ✓ Process: address space, resources, other general process attributes (e.g., privileges)
 - ✓ Execution state: PC, SP, registers, etc.
 - ✓ This execution state is usually called
 - a thread of control,
 - a thread, or
 - a lightweight process (LWP)

Thread Concept: Key Idea



Single and Multithreaded Processes




■ Single-threaded process

```
void func1(void *p) { ... }  
void func2(void *p) { ... }  
  
main()  
{  
    func1(...);  
    func2(...);  
    ...  
}
```

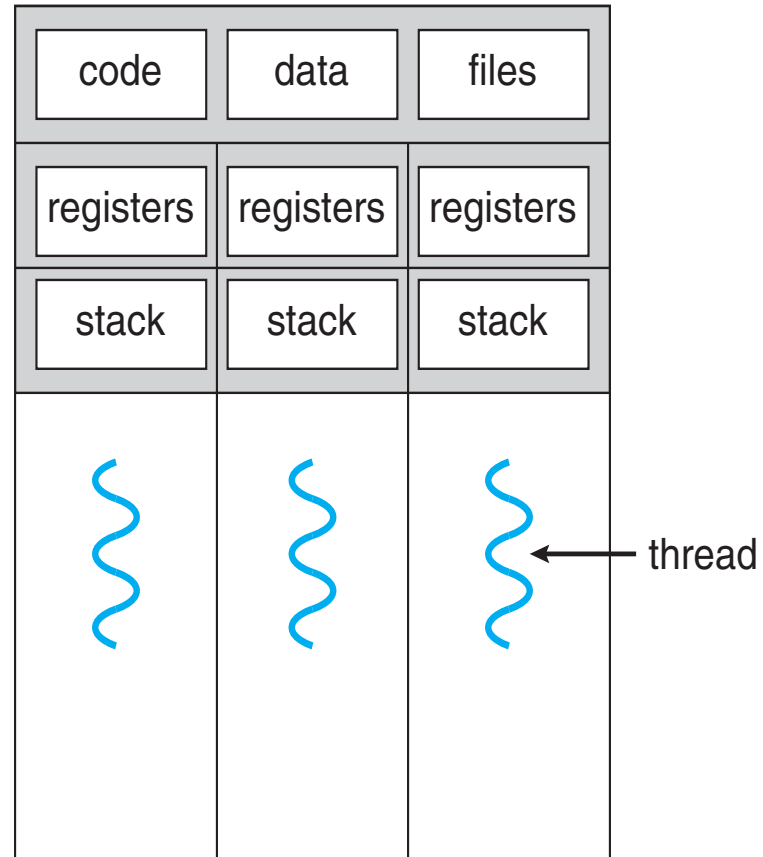
■ Multithreaded process

```
void func1(void *p) { ... }  
void func2(void *p) { ... }  
  
main()  
{  
    thread_create(func1, ...);  
    thread_create(func2, ...);  
    ...  
}
```

Single and Multithreaded Processes

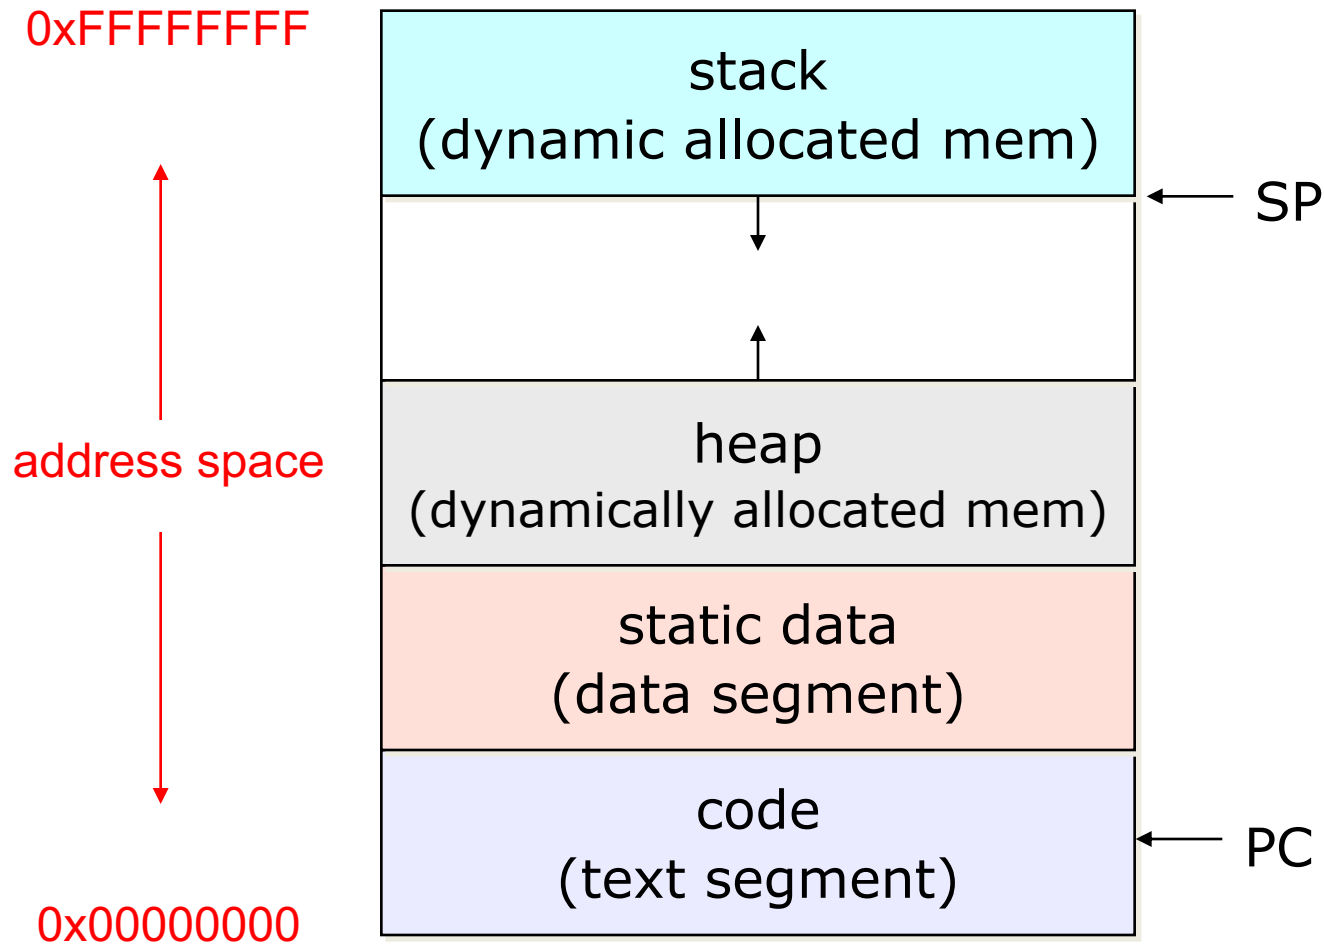


single-threaded process

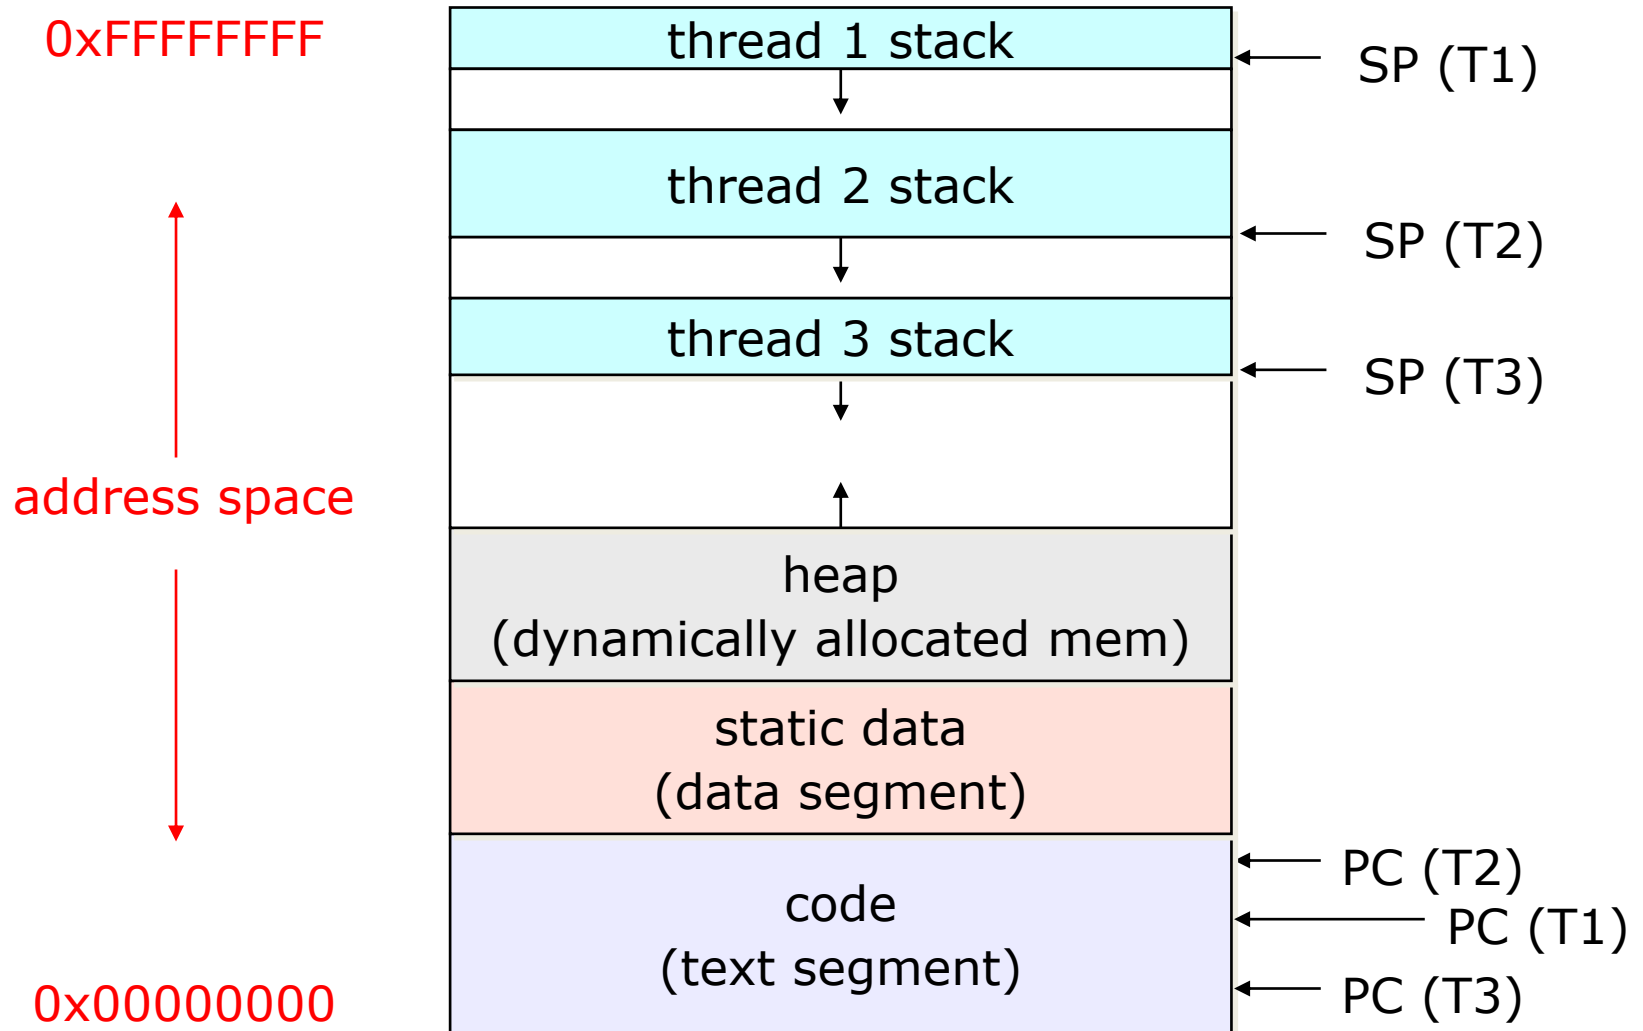


multithreaded process

Revisited: Process Address Space



Address Space with Threads



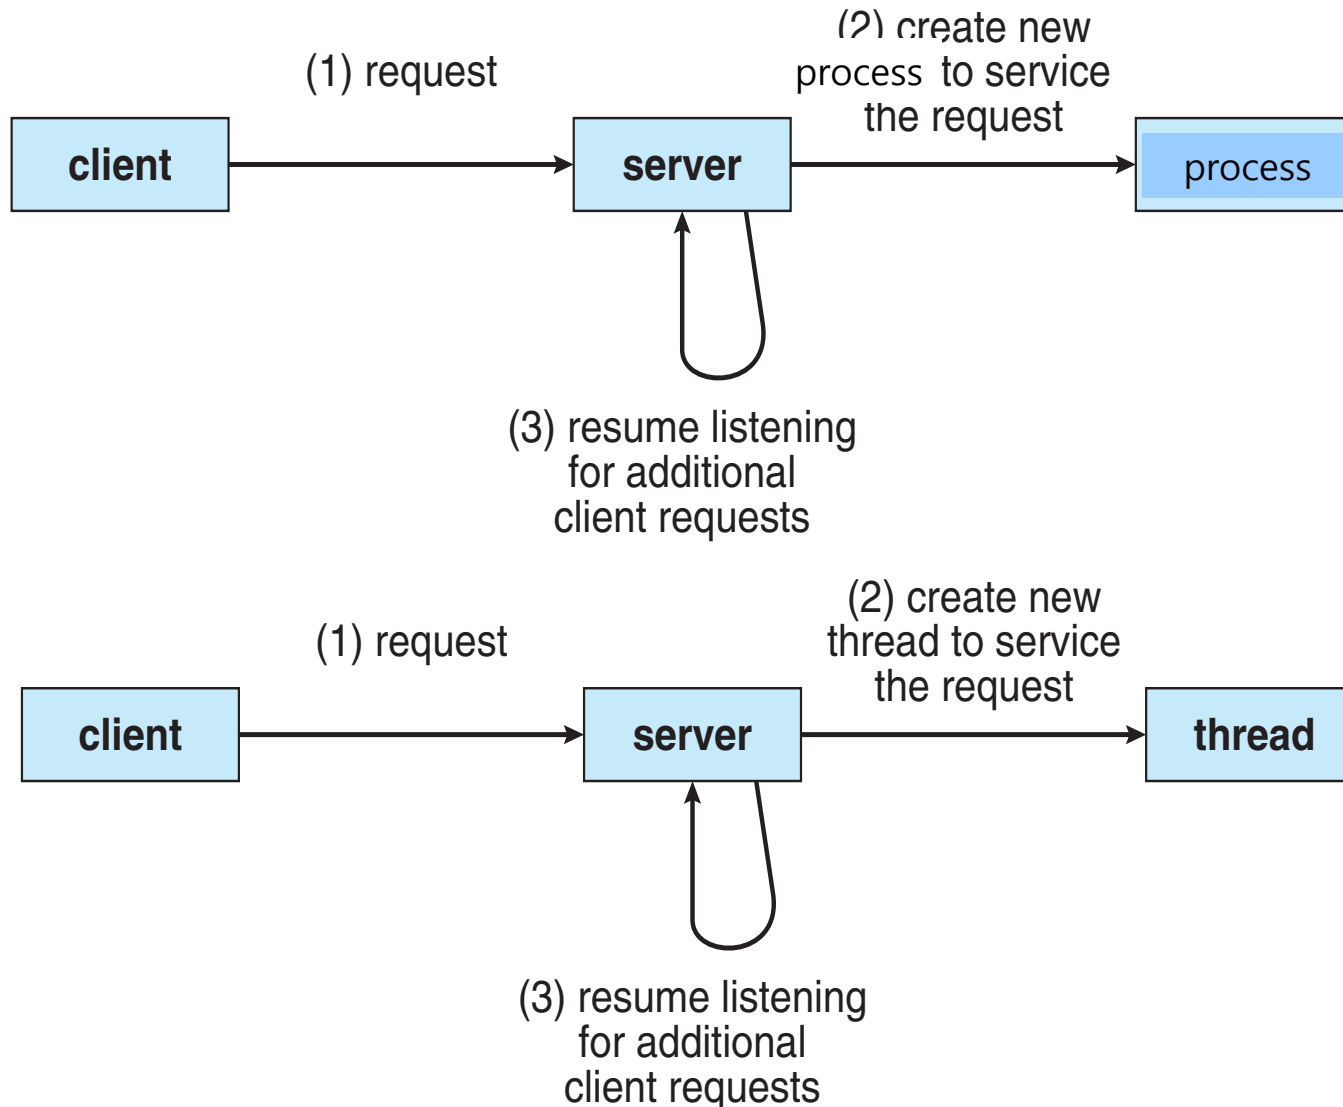
Concurrent Servers: Multiprocess Model

■ Web server example

- ✓ Using fork() to create new processes to handle requests in parallel is overkill for such a simple task

```
While (1) {  
    int sock = accept();  
    if ((pid = fork()) == 0) {  
        /* Handle client request */  
    } else {  
        /* Close socket */  
    }  
}
```

Concurrent Servers: Multiprocess → Multithread



Concurrent Servers: Multithread Model



- Using threads
 - ✓ We can create a new thread for each request

```
webserver ()
{
    While (1) {
        int sock = accept();
        thread_fork (handle_request, sock);
    }
}

handle_request (int sock)
{
    /* Process request */
    close (sock);
}
```

Single-Process (Iteration)

```
#include <stdio.h>
#define MAX_CMD 256

void DoCmd(char *cmd)
{
    printf("New command: %s\n", cmd);
    sleep(1);
    printf("Done\n");
}

int main()
{
    char cmd[MAX_CMD];

    while (1) {
        printf("CMD> ");
        fgets(cmd, MAX_CMD, stdin);
        if (cmd[0] == 'q')
            break;

        DoCmd(cmd);
    }

    return 0;
}
```

Multi-Processes



```
#include <stdio.h>
#define MAX_CMD 256

void DoCmd(char *cmd)
{
    printf("New command: %s\n", cmd);
    sleep(1); printf("Done\n");
    exit(0);
}

int main()
{
    char cmd[MAX_CMD]; int pid;
    while (1) {
        printf("CMD> ");
        fgets(cmd, MAX_CMD, stdin);
        if (cmd[0] == 'q') break;
        if ((pid = fork()) == 0) {
            DoCmd(cmd);
        }
    }
    #if 1
        else { wait(pid); }
    #endif
    return 0;
}
```

Multi-Threads



```
#include <stdio.h>
#include <pthread.h>
#define MAX_CMD 256

void DoCmd(char *cmd)
{
    printf("New command: %s\n", cmd);
    sleep(1); printf("Done\n");
    pthread_exit(NULL);
}

int main()
{
    char cmd[MAX_CMD]; pthread_t tid;
    while (1) {
        printf("CMD> ");
        fgets(cmd, MAX_CMD, stdin);
        if (cmd[0] == 'q') break;
        pthread_create(&tid, NULL, (void *)DoCmd, (void *)cmd);
    }
    pthread_join(tid, NULL);
    return 0;
}
```

Makefile



```
CC = gcc
CFLAGS =
LDFLAGS =
LIB = -lpthread
OBJ1 = iteration.o
OBJ2 = process.o
OBJ3 = thread.o
TARGET = iteration process thread

.SUFFIXES: .c .o
.c.o:
    $(CC) $(CFLAGS) -c $<

default: $(TARGET)
iteration: $(OBJ1)
    $(CC) -o $@ $(LDFLAGS) $(OBJ1)
process: $(OBJ2)
    $(CC) -o $@ $(LDFLAGS) $(OBJ2)
thread: $(OBJ3)
    $(CC) -o $@ $(LDFLAGS) $(OBJ3) $(LIB)
clean:
    rm -f *.o $(TARGET)
```

Two Programs on Arduino

```
#define LED 5
void setup() {
    pinMode(LED, OUTPUT);
}
void loop() {
    digitalWrite(LED, HIGH);
    delay(500);
    digitalWrite(LED, LOW);
    delay(500);
}

-----

#define BUZZER 6
enum { DO=262, RE=294, MI=330, FA=349 };
int Num = 4;
int Frequency[] = { DO, RE, MI, FA };
int Delay[] = { 300, 300, 300, 300 };
void setup() {
    pinMode(BUZZER, OUTPUT);
}
void loop() {
    for(int i=0; i<Num; i++) {
        tone(BUZZER, Frequency[i]);
        delay(Delay[i]);
    }
}
```


Single-Task on Arduino



```
#define LED      5
#define BUZZER  6
enum { DO=262, RE=294, MI=330, FA=349 };

int Num = 4;
int Frequency[] = { DO, RE, MI, FA };
int Delay[] = { 300, 300, 300, 300 };

void setup() {
    pinMode(LED, OUTPUT);
    pinMode(BUZZER, OUTPUT);
}

void loop() {

}
```

Multi-Tasks on Arduino



```
void LedTask() {
    while (1) {
        digitalWrite(LED, HIGH);
        delay(500);
        digitalWrite(LED, LOW);
        delay(500);
    }
}

void BuzzerTask() {
    while (1) {
        for(int i=0; i<Num; i++) {
            tone(BUZZER, Frequency[i]);
            delay(Delay[i]);
        }
    }
}


void setup() {
    xTaskCreate(LedTask, NULL, 200, NULL, 1, NULL);
    xTaskCreate(BuzzerTask, NULL, 200, NULL, 2, NULL);
    vTaskStartScheduler();
}

void loop() {
}
```


Multicore Programming



- Concurrent execution on a single-core system




- Parallel execution on a multicore system



Multicore Programming

■ Data vs. Task parallelism



Parallel Programming




- Pthreads (POSIX threads)
- OpenMP (Open Multi-Processing)
- Open MPI (Message Passing Interface)
- SIMD (Single Instruction Multiple Data)
- GPGPU (General Purpose computing on GPUs)
 - ✓ CUDA (Compute Unified Device Architecture)
 - ✓ OpenCL (Open Computing Language)

User Threads

- Thread management done by user-level threads library

- Example

- ✓ POSIX *Pthreads*
- ✓ Mach *C-thread*
- ✓ Solaris *threads*




Kernel Threads

- Supported by the Kernel

- ✓ Thread creation and management requires system calls

- Example

- ✓ Windows 95/98/NT/2000
- ✓ Solaris
- ✓ Tru64 UNIX
- ✓ BeOS
- ✓ Linux



User-level Thread vs. Kernel-level Threads

■ User-level threads


- ✓ The user-level threads library implements thread operations
- ✓ They are small and fast
- ✓ User-level threads are invisible to the OS
- ✓ OS may make poor decisions
 - E.g. blocking I/O
- ✓ Thread scheduling
 - Non-preemptive scheduling: `yield()`
 - Preemptive scheduling: timer through signal

■ Kernel-level threads


- ✓ All threads operations are implemented in the kernel
- ✓ The OS schedules all of the threads in a system
- ✓ Kernel threads are cheaper than processes
- ✓ They can still be too expensive

Multithreading Models

- Many-to-One





- One-to-One



- Many-to-Many

- Two-level

- ✓ Many-to-Many + One-to-One



Threading Issues



- Semantics of fork() and exec() **system call**
 - ✓ Two versions of fork()
- Thread cancellation
 - ✓ Asynchronous cancellation
 - ✓ Deferred cancellation
- Signal handling
 - ✓ To the thread to which the signal applies
 - ✓ To every thread in the process
 - ✓ To certain threads in the process
 - ✓ Assign a specific thread to receive all signals for the process
- Thread pools
 - ✓ Create a number of threads at process startup
- Thread specific data

Pthreads (POSIX threads)



■ Thread creation/termination

```
int pthread_create (pthread_t *tid,  
                  pthread_attr_t *attr,  
                  void *(start_routine)(void *),  
                  void *arg);
```

```
void pthread_exit (void *retval);
```

```
int pthread_join (pthread_t tid,  
                 void **thread_return);
```

Pthreads



■ Mutexes

```
int pthread_mutex_init  
    (pthread_mutex_t *mutex,  
     const pthread_mutexattr_t *mattr);
```

```
int pthread_mutex_destroy  
    (pthread_mutex_t *mutex);
```

```
int pthread_mutex_lock  
    (pthread_mutex_t *mutex);
```

```
int pthread_mutex_unlock  
    (pthread_mutex_t *mutex);
```

Pthreads



■ Condition variables

```
int pthread_cond_init  
    (pthread_cond_t *cond,  
     const pthread_condattr_t *cattr);
```

```
int pthread_cond_destroy  
    (pthread_cond_t *cond);
```

```
int pthread_cond_wait  
    (pthread_cond_t *cond,  
     pthread_mutex_t *mutex);
```

```
int pthread_cond_signal  
    (pthread_cond_t *cond);
```

```
int pthread_cond_broadcast  
    (pthread_cond_t *cond);
```

Windows Threads



■ Thread creation/termination

```
HANDLE CreateThread (lpThreadAttributes, dwStackSize,  
                    lpStartAddress, lpParameter,  
                    dwCreationFlags, lpThreadId);
```

```
void ExitThread  (dwExitCode);
```

Java Threads

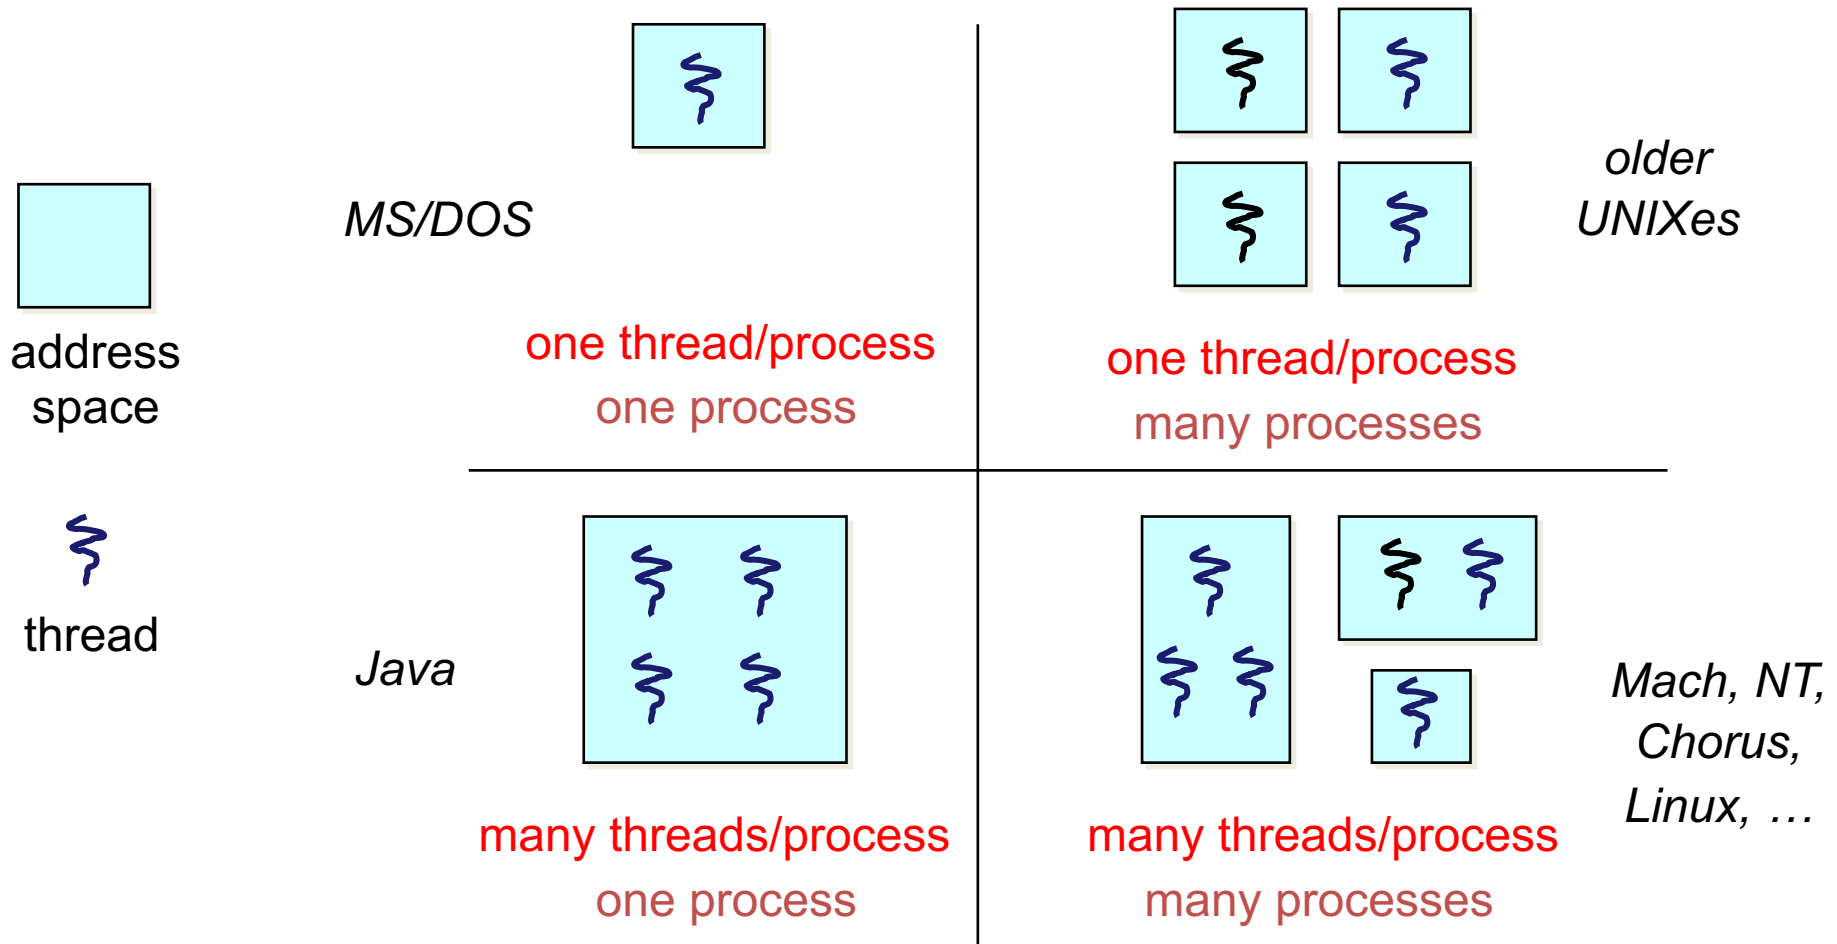


■ Thread creation/termination

Create a new class derived from `Thread` class
Override `run()` method

Create a new class that implements the `Runnable` interface

Threads Design Space



Thank You!
Q&A