

Operating System

Ch11: Mass-storage structure

BeomSeok Kim

Department of Computer Engineering

KyungHee University

passion0822@khu.ac.kr

Mass-Storage

- HDD (Hard Disk Drive)
- SSD (Solid State Disk)
- RAM disk
- Magnetic tape




Disk Attachment



- Disks may be attached one of two ways:
 1. **Host attached** via an I/O port
 2. **Network attached** via a network connection


Network-Attached Storage (NAS)

- Accessed via IP networks
- NFS, CIFS, etc.
- File-level access



Storage-Area Network (SAN)

- Accessed via private network dedicated for storages
- Use storage protocols such as SCSI or Fibre channel
- Block-level access
- File systems for SAN is another story (e.g. GFS)



Storage-Area Network (SAN)

- Storage array



Storage Architecture




		Protocols	
		Block	File
Interconnection	Non-IP Networks	IDE / SCSI (Direct) SAN (Storage Area Network: Fibre Channel)	DAFS (Direct Access File System: NFS over VIA)
	IP Networks	NBD (Network Block Device) iSCSI (SCSI over TCP/IP)	NAS (Network Attached Storage: NFS, CIFS)

■ Disks and the OS

- ✓ Disks are messy physical devices:
 - Errors, bad blocks, missed seeks, etc.
- ✓ The job of the OS is to hide this mess from higher-level software
 - Low-level device drivers (initiate a disk read, etc)
 - Higher-level abstractions (files, databases, etc.)
- ✓ The OS may provide different levels of disk access to different clients
 - Physical disk block (surface, cylinder, sector)
 - Disk logical block (disk block #)
 - Logical file (filename, block or record or byte #)

HDD Structure





■ Interacting with disks

- ✓ Specifying disk requests requires a lot of info:
 - Cylinder #, surface #, track #, sector #, transfer size, etc.
- ✓ Older disks required the OS to specify all of this
 - The OS needs to know all disk parameters
- ✓ Modern disks are more complicated
 - Not all sectors are the same size, sectors are remapped, etc.
- ✓ Current disks provide a higher-level interface (e.g. SCSI)
 - The disks exports its data as a logical array of blocks [0..N]
 - Disk maps logical blocks to cylinder/surface/track/sector
 - Only need to specify the logical block # to read/write
 - As a result, physical parameters are hidden from OS



■ Disk performance

✓ Performance depends on a number of steps


- **Seek**: moving the disk arm to the correct cylinder
 - depends on how fast disk arm can move (increasing very slowly)
- **Rotation**: waiting for the sector to rotate under head
 - depends on rotation rate of disk (increasing, but slowly)
- **Transfer**: transferring data from surface into disk controller, sending it back to the host
 - depends on density of bytes on disk (increasing, and very quickly)

✓ Disk scheduling:

- Because seeks are so expensive, the OS tries to schedule disk requests that are queued waiting for the disk
- Minimize seek time
- Seek time \approx seek distance

FCFS


- Illustration shows total head movement of 640 cylinders



SSTF


■ Shortest Seek Time First

- ✓ selects the request with the minimum seek time from the current head position
- ✓ a form of SJF scheduling
- ✓ may cause starvation of some requests




SCAN

■ *elevator algorithm*




C-SCAN

- Provides a more uniform wait time than SCAN



C-LOOK

- LOOK a version of SCAN, C-LOOK a version of C-SCAN



Selecting a Disk-Scheduling Algorithm



- Performance depends on the number and types of requests
 - ✓ SSTF is common and has a natural appeal
 - ✓ SCAN and C-SCAN perform better for systems that place a heavy load on the disk
- In general, unless there are request queues, disk scheduling does not have much impact
 - ✓ Important for servers, less so for PCs
- Modern disks often do the disk scheduling themselves
 - ✓ Disks know their layout better than OS, can optimize better
 - ✓ Ignores, undoes any scheduling done by OS

Disk Controllers



■ Intelligent controllers

- ✓ Nowadays, most disk controllers are built around a small CPU and have many kilobytes of memory
- ✓ They run a program written by the controller manufacturer to process I/O requests from the CPU and satisfy them
- ✓ Intelligent features:
 - Read-ahead: the current track
 - Caching: frequently-used blocks
 - Request reordering: for seek and/or rotational optimality
 - Request retry on hardware failure
 - Bad block identification
 - Bad block remapping: onto spare blocks and/or tracks

Swap-Space Management



- Swap-space
 - ✓ Virtual memory uses disk space as an extension of main memory

- Swap-space
 - ✓ can be carved out of the normal file system (Windows), or,
 - ✓ more commonly, it can be in a separate disk partition (Linux)


- Swap-space management
 - ✓ 4.3BSD allocates swap space when process starts
 - holds *text segment* (the program) and *data segment*
 - ✓ Kernel uses *swap maps* to track swap-space use
 - ✓ Solaris 2 allocates swap space only when a page is forced out of physical memory, not when the virtual memory page is first created

RAID




- Redundant Array of Inexpensive Disks
 - ✓ A storage system, not a file system
- Motivations
 - ✓ Use small cheap disks as a cost-effective alternative to large, expensive disks (I = Inexpensive)
 - ✓ Provide higher reliability and higher data-transfer (I = Independent)
- Improving reliability via redundancy
 - ✓ Mirroring (shadowing)
 - ✓ Parity or error-correcting codes
- Improving performance via parallelism
 - ✓ Data striping: bit-level vs. block-level

Multiple Disks vs. RAID



0	0	0	0
1	1	1	1
2	2	2	2
...
99	99	99	99



0
1
2
...
399

RAID Levels



■ RAID 0

- ✓ Non-redundant striping (no reliability)
- ✓ Data is broken down into blocks and each block is striped across multiple disks
- ✓ I/O performance is greatly improved by spreading the I/O load across many channels and drives
- ✓ Typically used in data rate intensive applications (e.g. video editing)

0	1	2	3
A0	A1	A2	A3
A4	A5	A6	A7
A8	A9	A10	A11
...

RAID Levels



■ RAID 1

- ✓ Mirrored disks
- ✓ Twice the read transaction rate of single disks, same write transaction rate as single disks
- ✓ Expensive, highest disk overhead
- ✓ No rebuild is necessary in case of a disk failure

0	1	2	3	4	5	6	7
A0	B0	C0	D0	A0	B0	C0	D0
A1	B1	C1	D1	A1	B1	C1	D1
A2	B2	C2	D2	A2	B2	C2	D2
...

RAID Levels



■ RAID 2

- ✓ Memory-style error-correcting codes (ECC)
- ✓ Each data word has its Hamming Code ECC word recorded on the ECC disks
- ✓ On read, the ECC code verifies correct data or corrects single bit errors
- ✓ ECC is embedded in almost all modern disk drives (e.g. SCSI)

0	1	2	3	4	5	6
A0	A0	A0	A0	Px	Py	Pz
A1	A1	A1	A1	Px	Py	Pz
A2	A2	A2	A2	Px	Py	Pz
A3	A3	A3	A3	Px	Py	Pz
...

RAID Levels

■ RAID 3

- ✓ Bit-interleaved parity
- ✓ Stripe parity is generated on writes, recorded on the parity disk and checked on reads
- ✓ Less storage overhead than RAID 2
- ✓ Requires hardware support for efficient use

0	1	2	3	4
A0	A0	A0	A0	P
A1	A1	A1	A1	P
A2	A2	A2	A2	P
A3	A3	A3	A3	P
...

RAID Levels



■ RAID 4

- ✓ Block-interleaved parity
- ✓ Parity for same rank blocks is generated on writes, recorded on the parity disk and checked on reads
- ✓ Very good read performance (same as RAID 0)
- ✓ Writes, however, require parity data be updated each time
- ✓ No advantages over RAID 5 and does not support multiple simultaneous write operations

0	1	2	3	4
A0	A1	A2	A3	P
A4	A5	A6	A7	P
A8	A9	A10	A11	P
...

RAID Levels



■ RAID 5

- ✓ Block-interleaved distributed parity
- ✓ Parity for blocks in the same rank is generated on writes, recorded in a distributed location
- ✓ Can speed up small writes in multiprocessing systems, since the parity disk does not become a bottleneck

0	1	2	3	4
A0	A1	A2	A3	P
A4	A5	A6	P	A7
A8	A9	P	A10	A11
A12	P	A13	A14	A15
...

RAID Levels

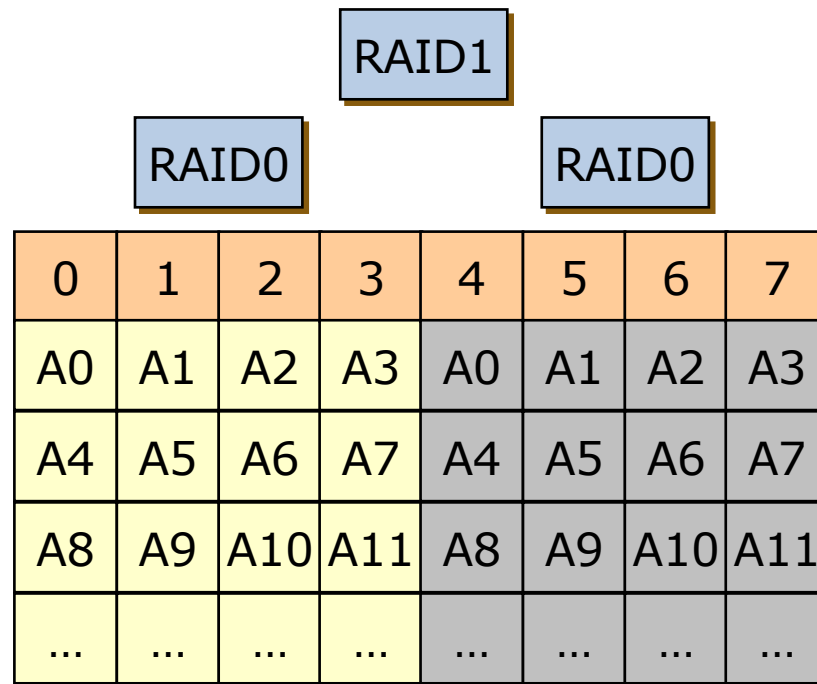
■ RAID 6

- ✓ P+Q redundancy scheme
- ✓ RAID 5 + extra redundancy information to guard against multiple disk failure
- ✓ Use error-correcting codes instead of parity

0	1	2	3	4	5
A0	A1	A2	A3	Px	Py
A4	A5	A6	Px	Py	A7
A8	A9	Px	Py	A10	A11
A12	Px	Py	A13	A14	A15
...

- RAID 0+1

- ✓ A set of disks are striped, and then the stripe is mirrored to another, equivalent stripe
- ✓ RAID 0 provides performance, while RAID 1 provides the reliability
- ✓ A single drive failure will cause the whole array to become, in essence, a RAID 0 array



Thank You!
Q&A